

Cloud and Big Data

Dr. Sambit Sahu, VP AI Foundations, Capital One



Course Objective

- **Graduate/Adv Undergrad level course on Cloud Computing**
 - Focus is on learning and building extremely large scale systems and applications leveraging Cloud.
 - Learn concepts as well as hands-on experience by using real cloud and cloud technologies.
 - Three key objectives: build applications using cloud, learn building blocks/services required to design large scale applications, computing in a cluster
- **We learn cloud technologies by using real clouds - Amazon AWS, Google Cloud, Hadoop & Spark platform.**
- **Required background**
 - Programming experience with one of the following Java/Python
 - Concepts of web services and applications
 - Optional: Operating Systems fundamentals, networking concepts

Three Main Components

- **Cloud Programming**

- Basic cloud concepts
- Amazon AWS cloud and services
- Google App Engine
- ***Build a large application leveraging cloud***

- **Cloud building blocks and services**

- Compute Cloud: Virtualization concepts,
- Storage Cloud, Cloud Database
- Message Queues
- Cloud devOps
- ***Design pattern in extreme scale backend engineering***

- **Big Data Platform and Programming**

- Hadoop eco-system, Map-Reduce, HDFS
- Spark with RDD and dataframes
- Intelligent systems and pipeline
- ***Web scale data computing and pipeline***

Tentative Schedule

Date	Lecture	Reading Papers	Assignment/HW
09/05	Intro to Cloud		
09/12	Building Applications using AWS Cloud	GFS	
09/19	Large Scale Systems Design Patterns	Big Table	A1 Release
09/26	Containers, Kubernetes and Micro-services	DynamoDB	
10/03	Cloud DevOps	Kafka	
10/10	Messages Queues (Kafka) and Streams	Borg	A2 release
10/17	Quiz 1	Map Reduce	
10/24	Cluster Computing with Hadoop	Spark RDD	
10/31	Cluster Computing with Spark	Spanner	
11/07	Spark Dataframes and Data Pipelines		A3 release
11/14	Database: SQL, noSQL, Elastic Search		
11/21	Quiz 2		

Course Structure

■ **Lecture Structure**

- Each lecture will have a theme topic. First 1 hour 30 minutes lecture and demonstration by the instructor.
- Last 20 minutes students to lead discussions from the reading paper lists.

■ **Assignments/Exercises**

- Reading list – consists of 10 landmark papers in the area of large scale systems (Google File System, Map Reduce, Spanner, Hadoop, Amazon DynamoDB, Kafka, Borg, Spark RDD etc.)
 - Submit paper summaries
 - Three Programming Assignments
 - Course Project: you can conduct this in a group of 3-4 students

■ **Communication Channel:** Slack, Brightspace

Grading and requirements

- Class participations (paper critics/quizzes) -- 5% grade
 - Paper discussions
 - Paper summaries
- Mid-Term covering concepts, design and coding -- 25%
- Assignments – 35% grade
 - 3 programming assignments stressed on technologies and programming
- Course project -- 35% grade
 - Students may team upto size of four
- Submission process – Brightspace

Project: Learn how to innovate in this space

- **Objective is to learn how to innovate in this space**

- **Four phases to your project**
 1. Concept and business idea
 2. Technology viability and architecture
 3. Execution planning and prototyping
 4. Demo, socialization and review

- **Few suggestion**
 - Form your team carefully – asking, interviewing your teammates. Float around some ideas,, kick the tire. Take a look at lot of recent startups that are bought by Google, Apple, FB, Amazon etc. **Take a look at beta.list**
 - I will provide a set of skeleton ideas you could choose from.

What you need to do soon

- Get account on few popular clouds
 - Amazon AWS (EC2, S3)
 - Use AWS coupons if first time user
 - Build a static website using S3:
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/HostingWebsiteOnS3Setup.html>
 - Create a Virtual Machine using AWS EC2
- Course Project
 - Substantial portion of your grade depends on final course project
 - I will provide sample project ideas/topics
 - You need to have a team and a project topic submitted by end of 3rd week

Reference Books

- Too many topics – so not any particular good book. So attend lectures, read my mind (hopefully the good side of it!), learn from using real cloud and code (a lot of it actually!), and Google!
- Reference books
 - AWS in Action
 - Learning Spark Programming
 - Kubernetes in Action

What is a WebApp?

Webapp: A web application (web app) is **an application program that is stored on a remote server and delivered over the internet through a browser interface**. Web services are web apps by definition and many, although not all, websites contain web apps

Examples:

- ????
- ????

Key Components:

- backend server stack
 - web server
 - application server
 - database
- frontend/client
- APIs
- network connectivity

Web Application Primer

Different Components of a WebApp

Client and Server Architecture

Client (frontend)

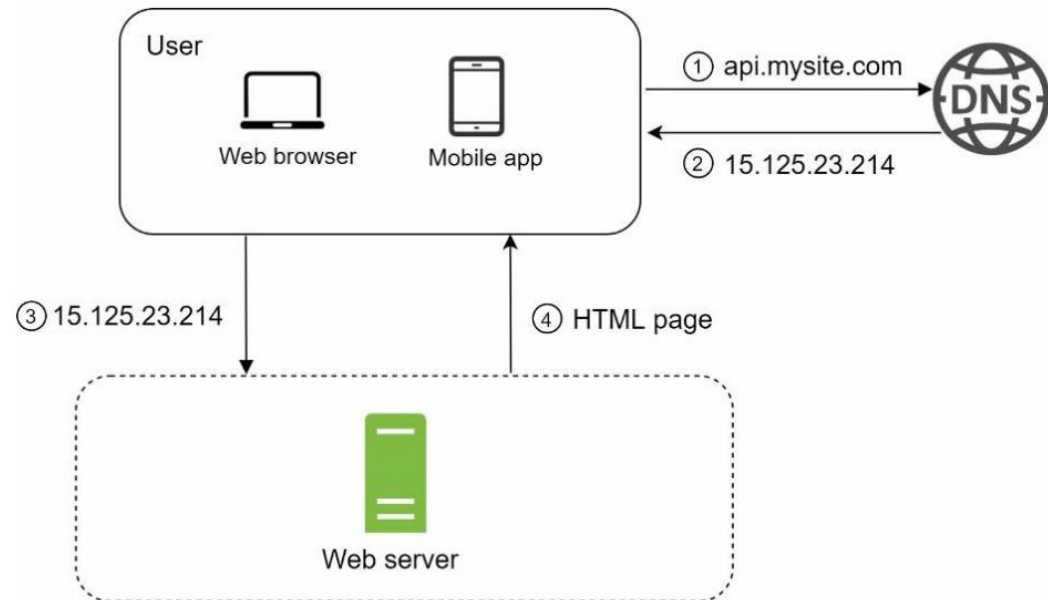
- webapps, mobile apps
- React (FB), Angular 4 etc

Server (backend)

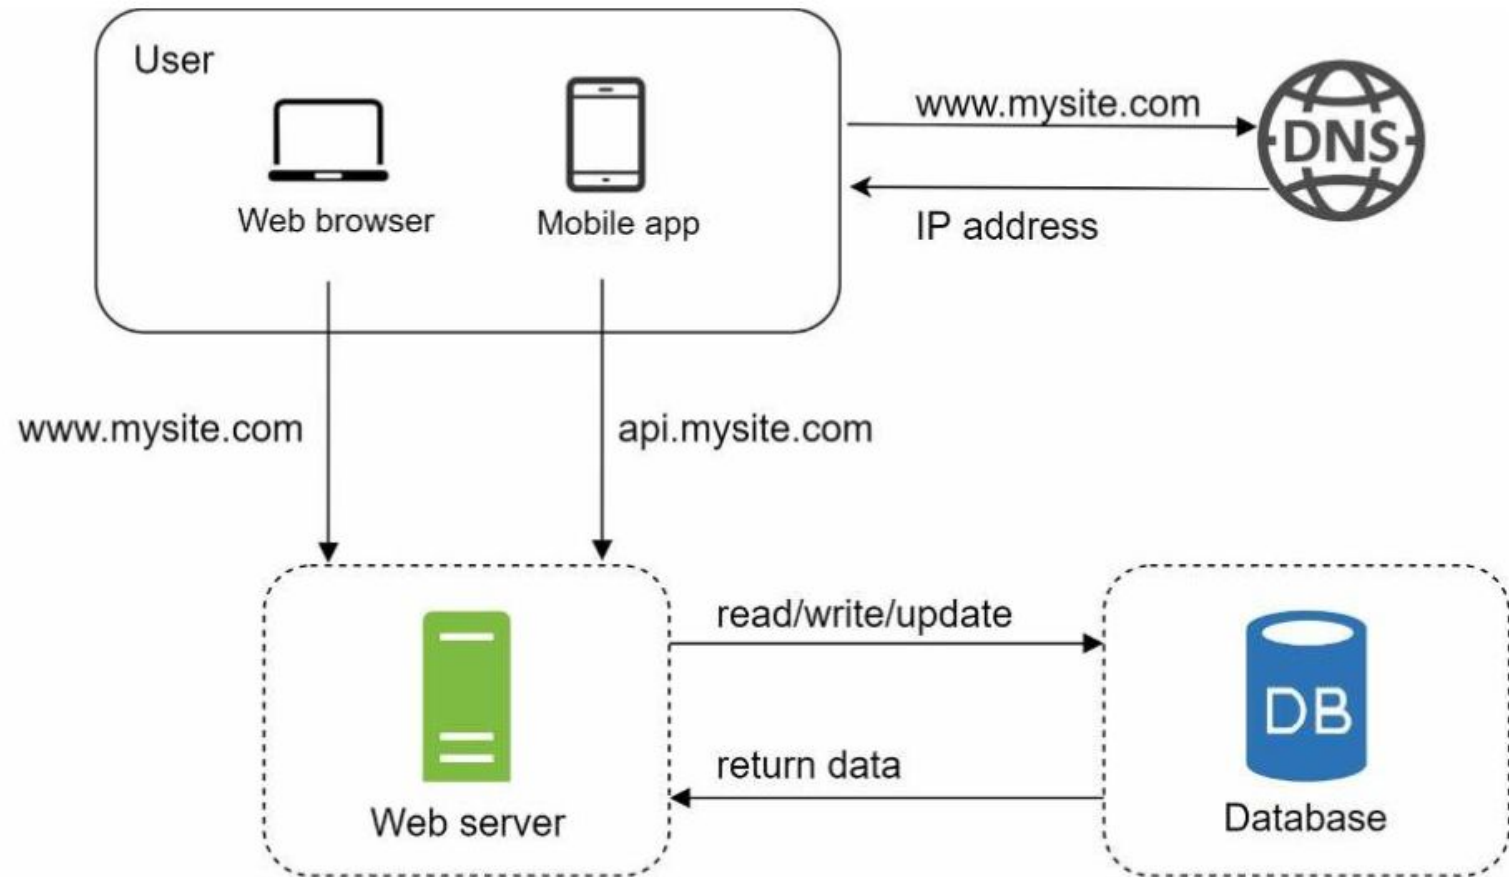
- three tier backend platform
- web server: node.js etc

Web APIs

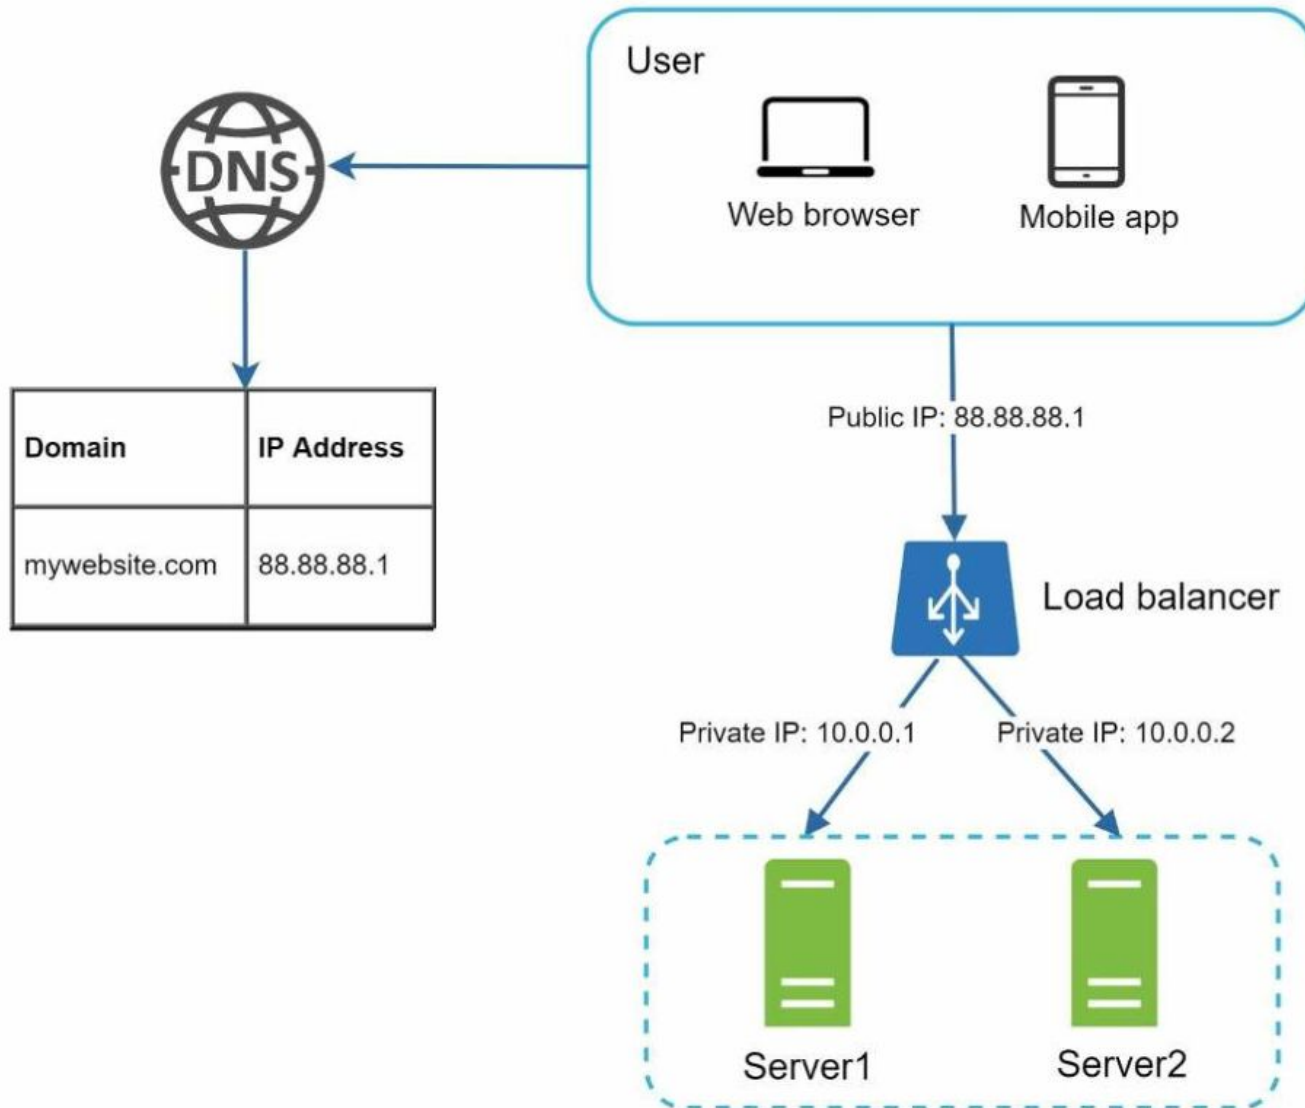
- frontend requests a services running in the backend via a HTML/web request
- REST based APIs



A Stateful WebApp



Scaling the backend of Webapp



Scaling the Database

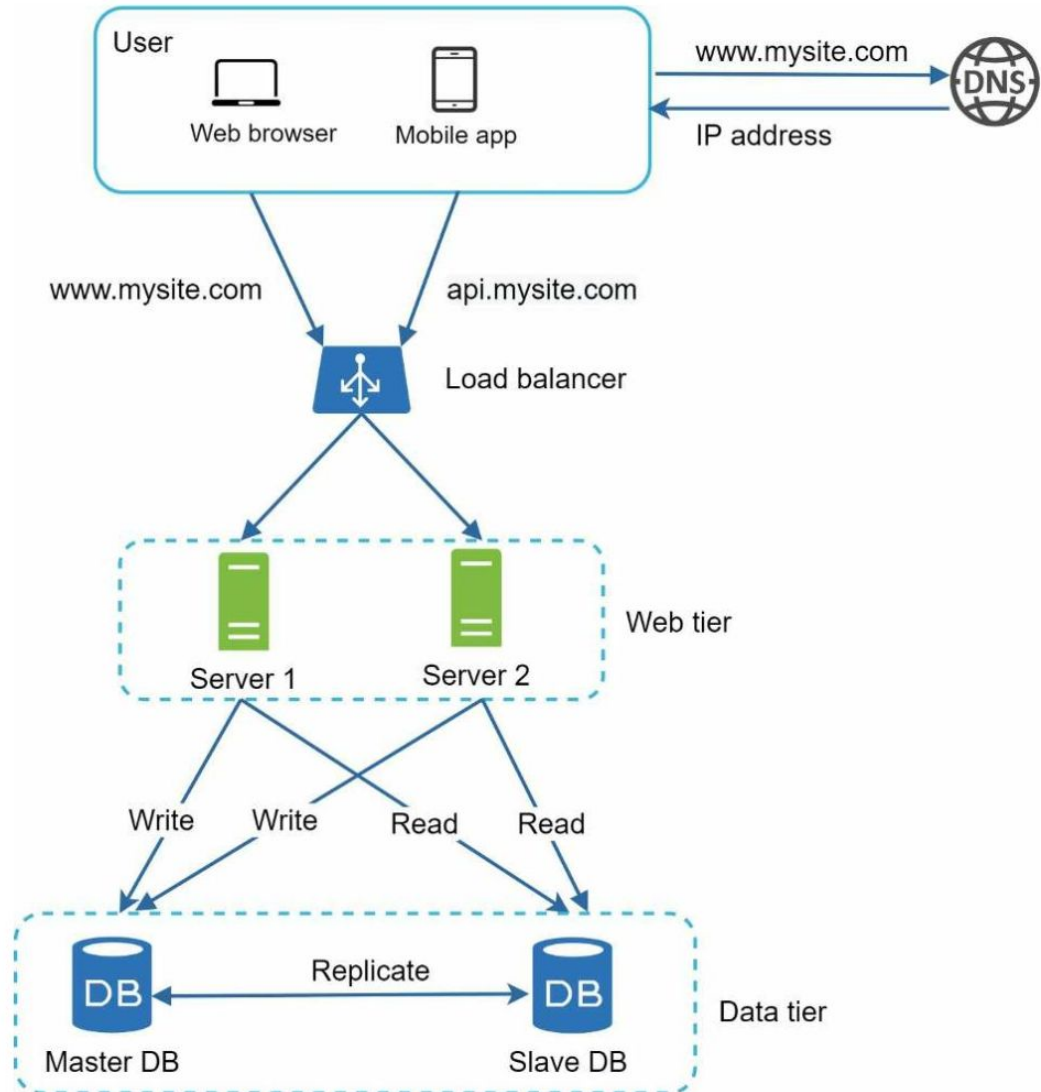
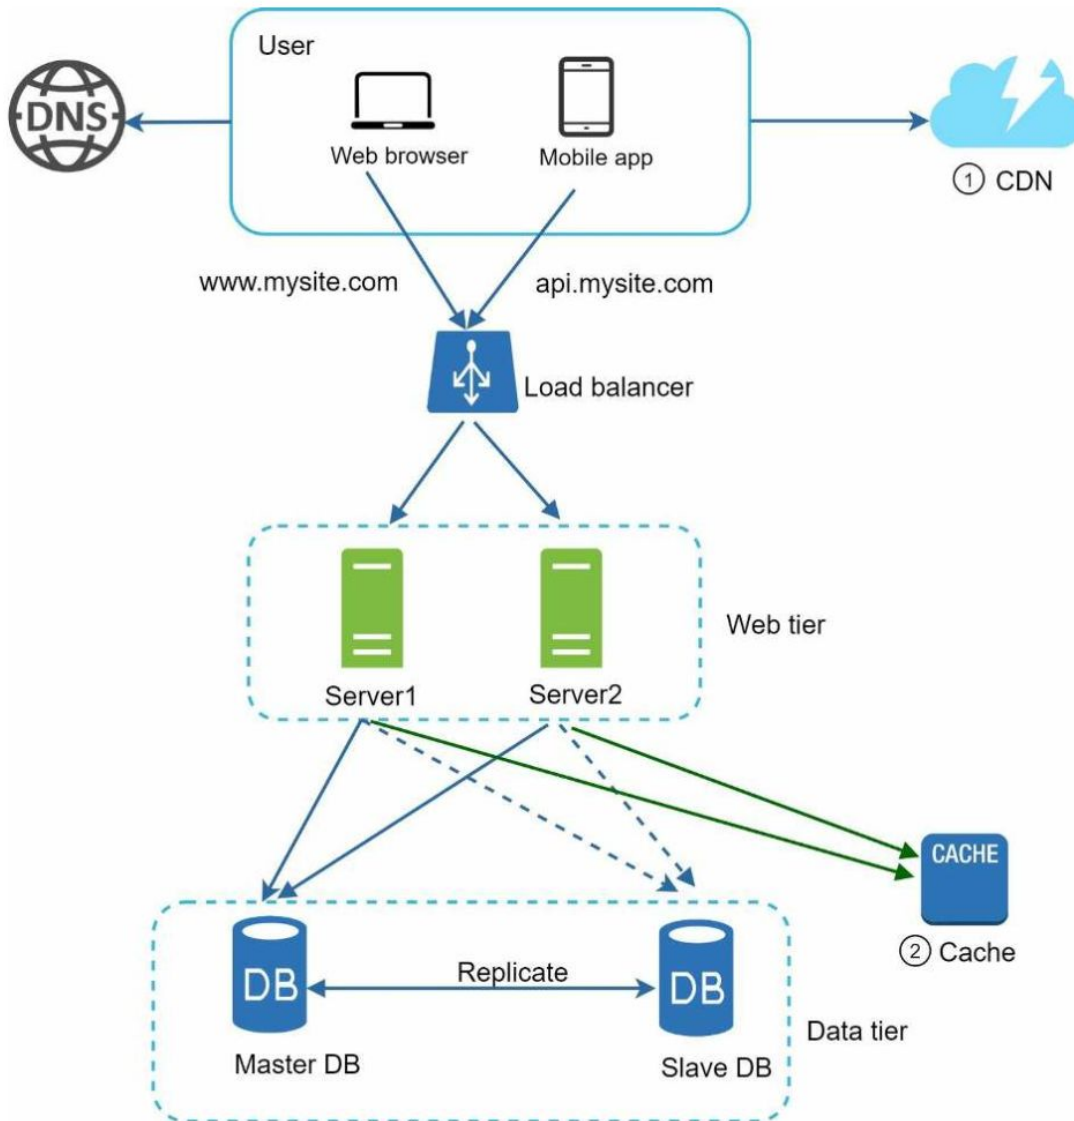


Figure 4.6

Scaling the Performance



What is Cloud?

- Allows users to request computing/storage resources and services through web interfaces
- You do not need to own or install or manage these resources.
- Pay as you go - Resources on-demand
- Elastic: Use as much as you want or as less as you want
 - Users can assume infinite amount of compute and storage resources are available.
 - Users can request resources when and what they need and release/remove resources when they don't need.
- Compute and storage resources are now treated as software entities. You get access to such resources programmatically – not by physical hardware anymore!
- Example Clouds:
 - AWS: aws.amazon.com
 - Google Cloud: <https://console.cloud.google.com/>
 - Azure: <https://azure.microsoft.com/en-us/products/>

Why Cloud?

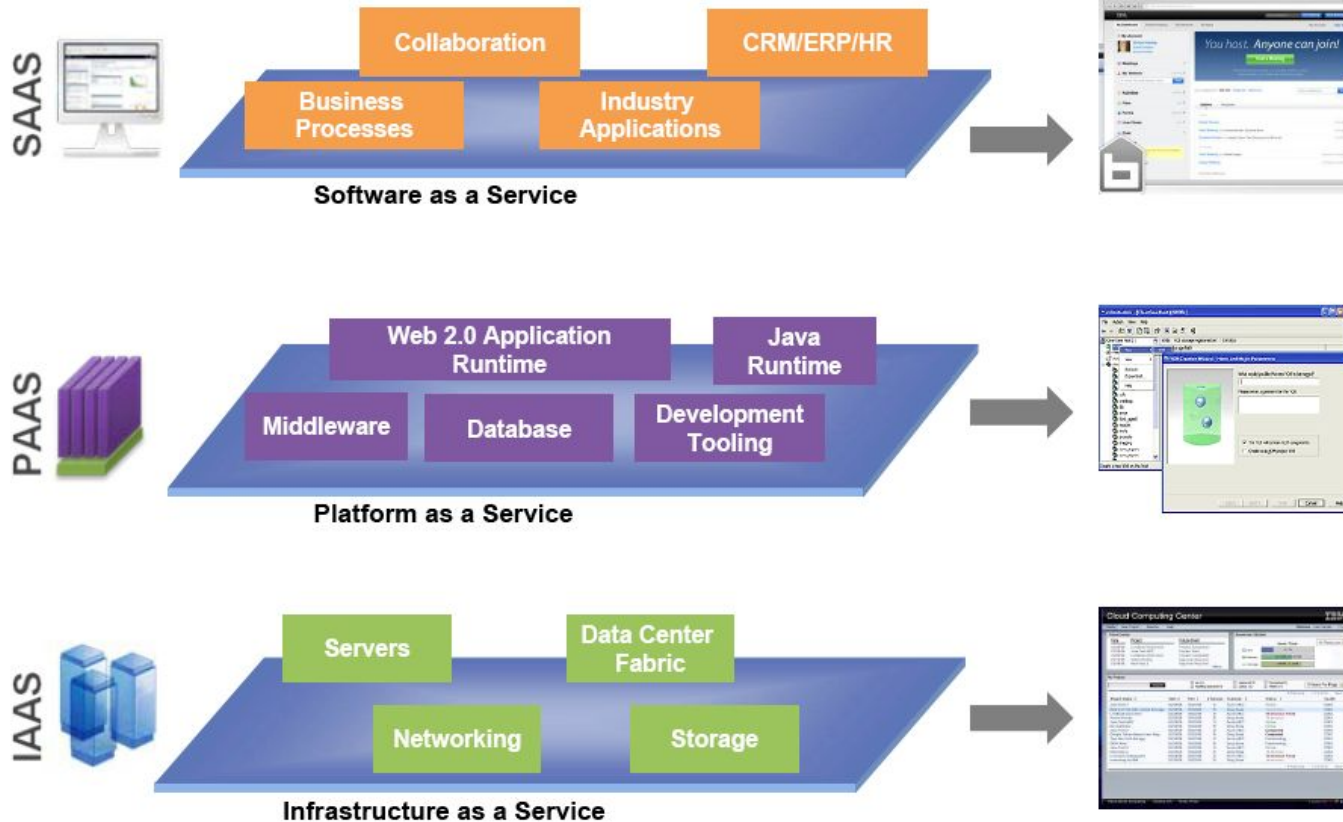
- You can get as many as 1000 machines for an hour for a few dollars to run a complex application!
- You don't need to manage, maintain or fix any machines!
- You can use as little as 1 machine or as many as 10000 machines depending on what your current needs are!
- Two key focus: on-demand and elastic!

Service Models

- *Cloud Software as a Service (SaaS)*. The capability provided to the consumer is to use the provider's applications running on a cloud infrastructure. The applications are accessible from various client devices through a thin client interface such as a web browser (e.g., web-based email). The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, storage, or even individual application capabilities, with the possible exception of limited user-specific application configuration settings.
- *Cloud Platform as a Service (PaaS)*. The capability provided to the consumer is to deploy onto the cloud infrastructure consumer-created or acquired applications created using programming languages and tools supported by the provider. The consumer does not manage or control the underlying cloud infrastructure including network, servers, operating systems, or storage, but has control over the deployed applications and possibly application hosting environment configurations.
- *Cloud Infrastructure as a Service (IaaS)*. The capability provided to the consumer is to provision processing, storage, networks, and other fundamental computing resources where the consumer is able to deploy and run arbitrary software, which can include operating systems and applications. The consumer does not manage or control the underlying cloud infrastructure but has control over operating systems, storage, deployed applications, and possibly limited control of select networking components (e.g., host firewalls).

Different Cloud Offerings: A Layered Perspective

The Layers of IT-as-a-Service



- Higher the stack, less control but more automation for user
- Lower the stack, more control but more responsibility for user

Cloud Computing Delivery Models

Flexible Delivery Models

Public ...

- Service provider owned and managed
- Access by subscription
- Delivers select set of standardized business process, application and/or infrastructure services on a flexible price per use basis

Cloud
Services

Cloud
Computing
Model

Private ...

- Privately owned and managed.
- Access limited to client and its partner network.
- Drives efficiency, standardization and best practices while retaining greater customization and control

Hybrid ...

Access to client, partner network, and third party

....Standardization, capital preservation, flexibility and time to deploy

.... Customization, efficiency, availability, resiliency, security and privacy


ORGANIZATION → CULTURE → GOVERNANCE

...service sourcing and service value

Cloud Computing Economics

- Three useful usage scenarios
 - Load varying with time
 - Demand unknown in advance
 - Batch analytics that can benefit from huge number of resources for a short time duration
- Why pay-as-you-go model makes sense economically even if costs higher than buying a server and depreciating the h/w
 - Extreme elasticity
 - Transference of risk (of over provisioning)

Figure 2. (a) Even if peak load can be correctly anticipated, without elasticity we waste resources (shaded area) during nonpeak times. (b) Underprovisioning case 1: potential revenue from users not served (shaded area) is sacrificed. (c) Underprovisioning case 2: some users desert the site permanently after experiencing poor service; this attrition and possible negative press result in a permanent loss of a portion of the revenue stream.



Let's use a IaaS Cloud (Amazon EC2)

- <http://aws.amazon.com/console/>
- Amazon EC2 console based provisioning demo
- Build a static website:
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/HostingWebsiteOnS3Setup.html>

Amazon AWS console (EC2 view)

The screenshot displays the Amazon AWS console interface, specifically the Amazon EC2 Console Dashboard. The top navigation bar includes the AWS logo, the user's name (sambit sahu), and links for Settings and Sign Out. The left sidebar contains a navigation menu with categories like INSTANCES, IMAGES, ELASTIC BLOCK STORE, and NETWORKING & SECURITY. The main content area is titled 'Amazon EC2 Console Dashboard' and features several sections: 'Getting Started' with a 'Launch Instance' button, 'My Resources' showing a summary of EC2 resources (1 Running Instance, 1 EBS Volume, 2 Key Pairs, 0 Load Balancers, 0 Elastic IPs, 0 EBS Snapshots, 6 Security Groups, 0 Placement Groups), and 'Service Health' showing the current status of Amazon EC2 (US East - N. Virginia) as 'Service is operating normally'.

Navigation

Region: US East

EC2 Dashboard

INSTANCES

- Instances
- Spot Requests

IMAGES

- AMIs
- Bundle Tasks

ELASTIC BLOCK STORE

- Volumes
- Snapshots

NETWORKING & SECURITY

- Elastic IPs
- Security Groups
- Placement Groups
- Load Balancers
- Key Pairs

Amazon EC2 Console Dashboard

Getting Started

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

[Launch Instance](#)

Note: Your instances will launch in the US East (Virginia) region.

My Resources

You are using the following Amazon EC2 resources in the US East (Virginia) region: [Refresh](#)

- 1 Running Instance
- 0 Elastic IPs
- 1 EBS Volume
- 0 EBS Snapshots
- 2 Key Pairs
- 6 Security Groups
- 0 Load Balancers
- 0 Placement Groups

Service Health

Current Status	Details
	Amazon EC2 (US East - N. Virginia) Service is operating normally

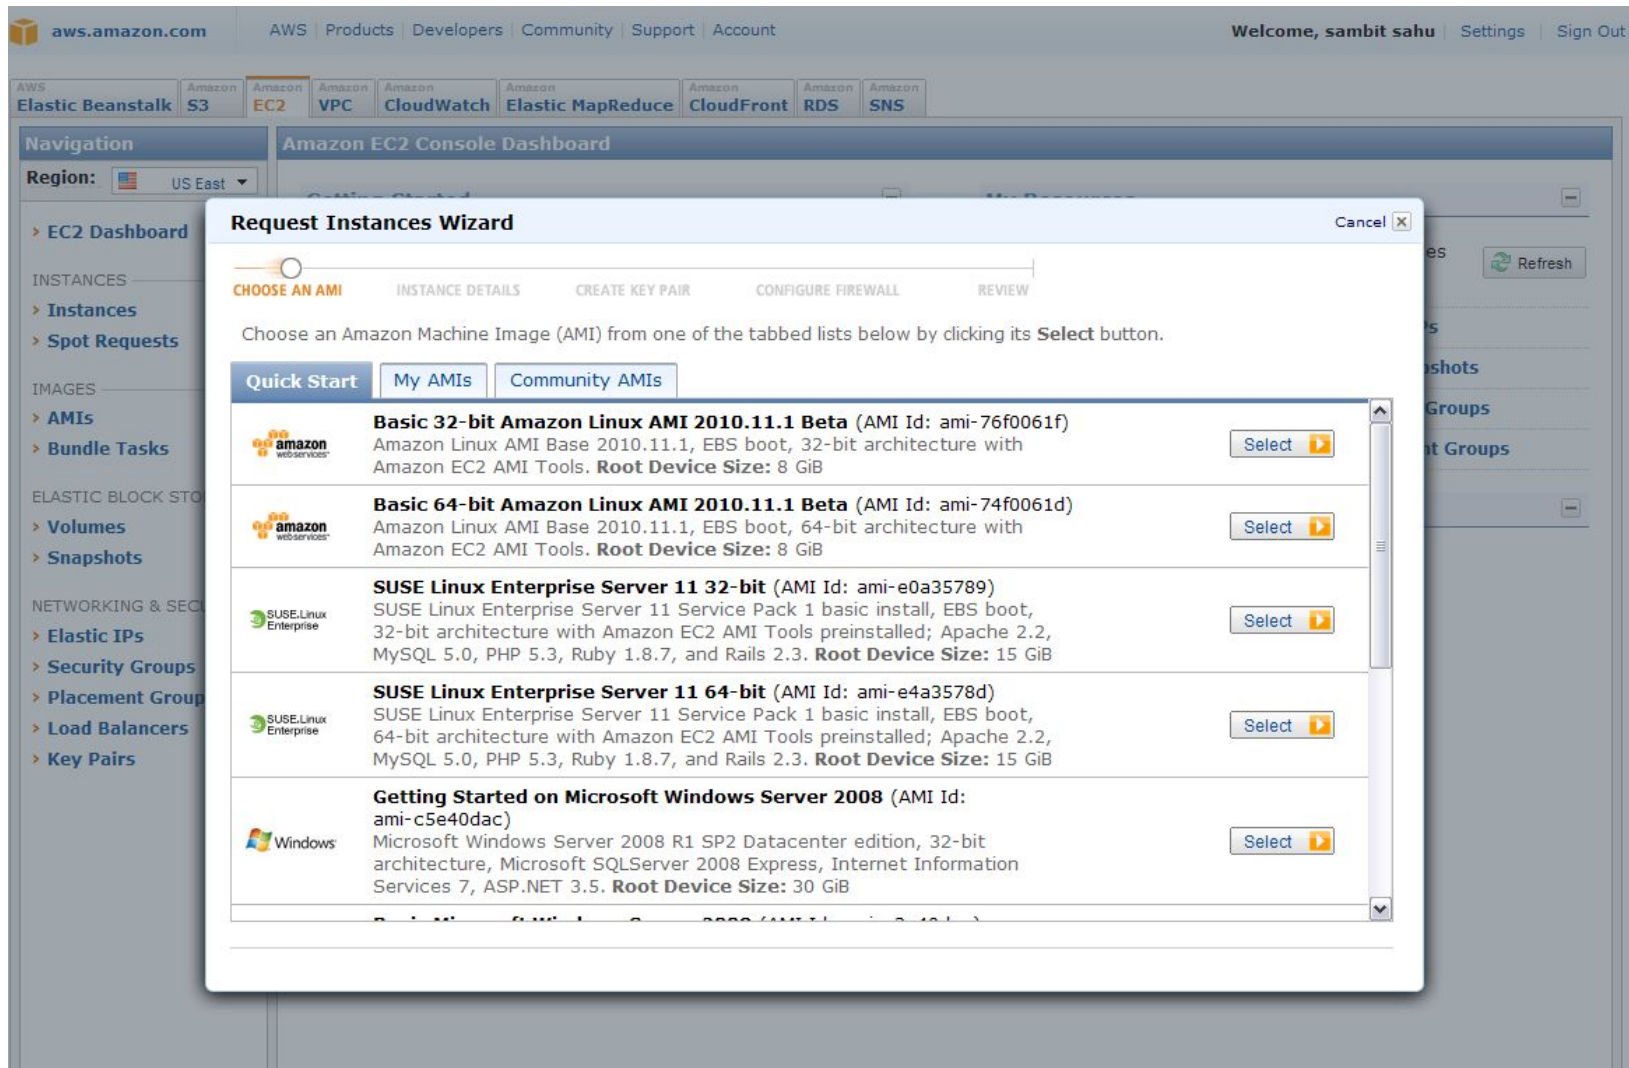
[View complete service health details](#)

Related Links

- Documentation
- All EC2 Resources
- Forums
- Feedback
- Report an Issue

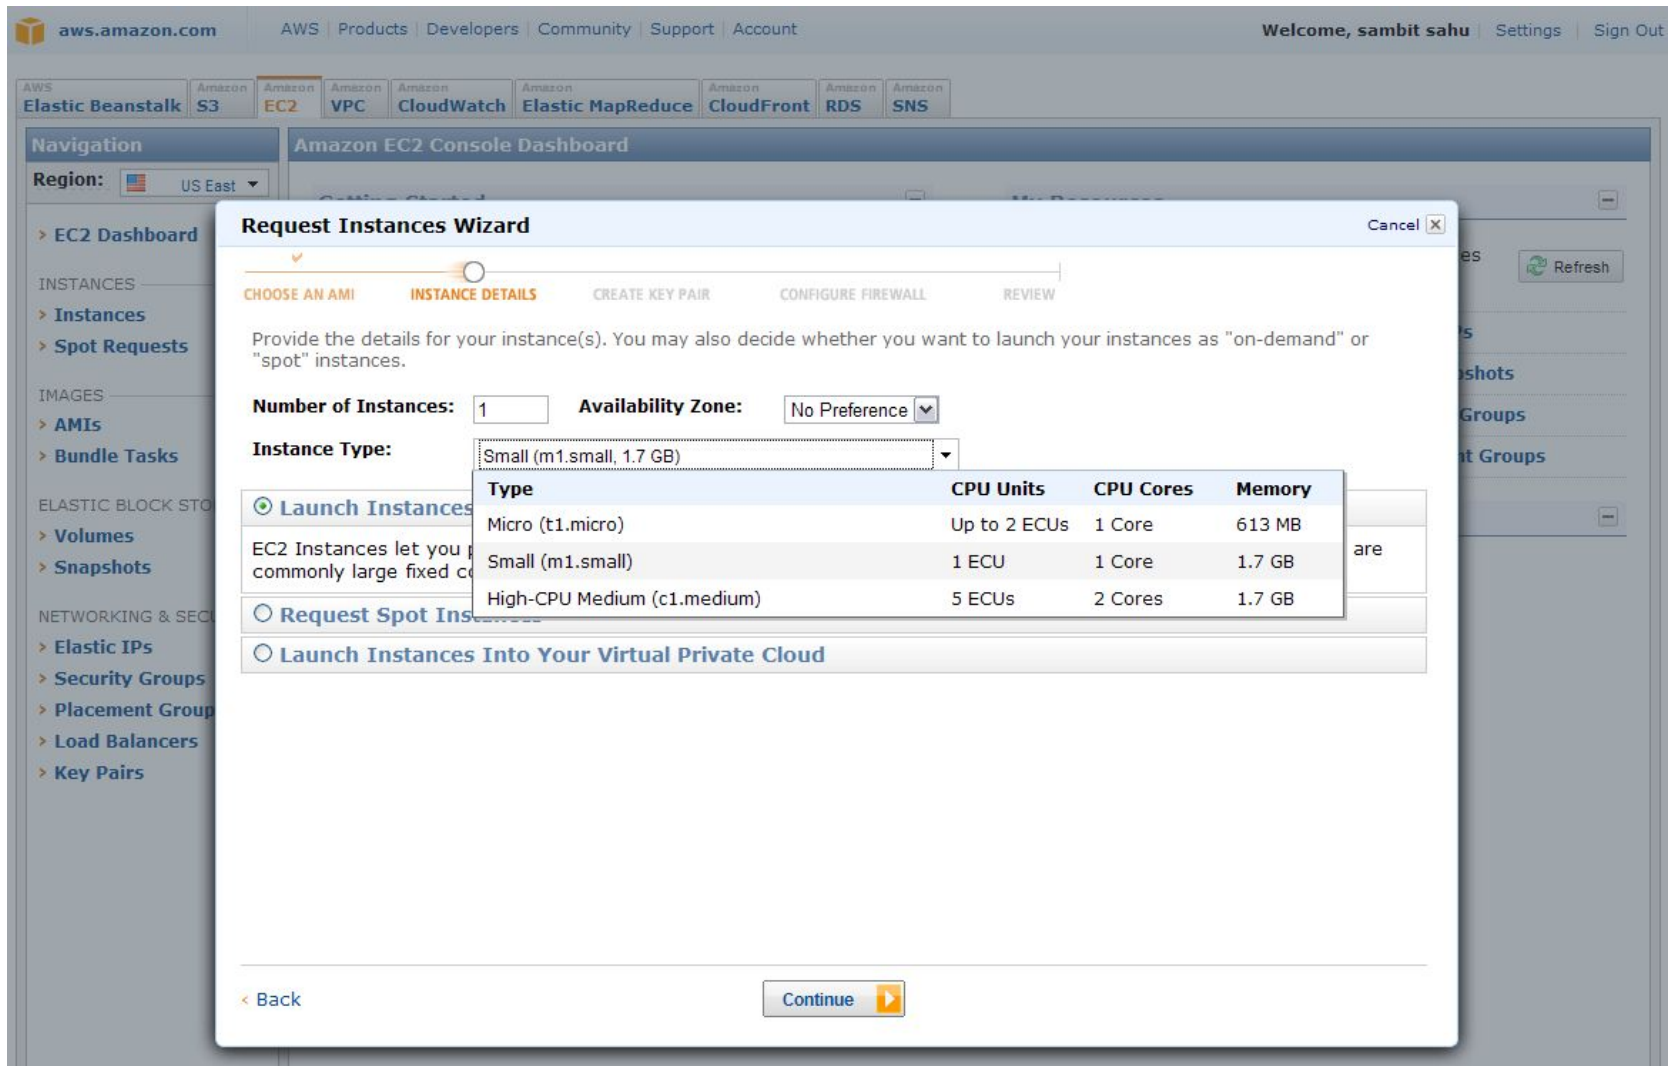
- User logs in with AWS credentials

User launches request instance □ a list of prebuilt stack is provided



- AWS shows a list of available pre-built base software stack (called **Virtual Appliances**) user may request to add to the machine

User can choose the resource size (CPU, mem choices)



Request Instances Wizard

CHOOSE AN AMI **INSTANCE DETAILS** CREATE KEY PAIR CONFIGURE FIREWALL REVIEW

Provide the details for your instance(s). You may also decide whether you want to launch your instances as "on-demand" or "spot" instances.

Number of Instances: 1 **Availability Zone:** No Preference

Instance Type: Small (m1.small, 1.7 GB)

Type	CPU Units	CPU Cores	Memory
Micro (t1.micro)	Up to 2 ECUs	1 Core	613 MB
Small (m1.small)	1 ECU	1 Core	1.7 GB
High-CPU Medium (c1.medium)	5 ECUs	2 Cores	1.7 GB

Launch Instances

EC2 Instances let you launch commonly large fixed capacity instances.

Request Spot Instances

Launch Instances Into Your Virtual Private Cloud

[Back](#) [Continue](#)

- Instance request wizard guides through resource choices

User specifies security/access configurations

The screenshot displays the AWS Management Console interface. At the top, the navigation bar includes the AWS logo, the URL 'aws.amazon.com', and links for 'Products', 'Developers', 'Community', 'Support', and 'Account'. The user is logged in as 'sambit sahu'. Below the navigation bar, a horizontal menu lists various AWS services: Elastic Beanstalk, S3, EC2 (highlighted), VPC, CloudWatch, Elastic MapReduce, CloudFront, CloudFormation, RDS, SNS, and IAM. The left sidebar contains a 'Navigation' menu with sections for 'Region' (set to US East (Virginia)), 'INSTANCES' (EC2 Dashboard, Instances, Spot Requests, Reserved Instances), 'IMAGES' (AMIs, Bundle Tasks), 'ELASTIC BLOCK STORE' (Volumes, Snapshots), and 'NETWORKING & SECURITY' (Security Groups, Elastic IPs, Placement Groups, Load Balancers, Key Pairs). The main content area shows the 'Amazon EC2 Console Dashboard' with tabs for 'Getting Started' and 'My Resources'. A 'Request Instances Wizard' dialog box is open in the center, featuring a progress bar with five steps: 'CHOOSE AN AMI', 'INSTANCE DETAILS', 'CREATE KEY PAIR' (the current step), 'CONFIGURE FIREWALL', and 'REVIEW'. The dialog box contains the following text: 'Public/private key pairs allow you to securely connect to your instance after it launches. To create a key pair, enter a name and click **Create & Download your Key Pair**. You will then be prompted to save the private key to your computer. Note, you only need to generate a key pair once - not each time you want to deploy an Amazon EC2 instance.' Below this text, there are three radio button options: 'Choose from your existing Key Pairs' (selected), 'Create a new Key Pair', and 'Proceed without a Key Pair'. The 'Choose from your existing Key Pairs' option has a dropdown menu showing 'IM2011'. At the bottom of the dialog box, there are '< Back' and 'Continue >' buttons.

aws.amazon.com | AWS | Products | Developers | Community | Support | Account | Welcome, sambit sahu | Settings | Sign Out

Navigation
Region: US East (Virginia)

Amazon EC2 Console Dashboard

Getting Started | My Resources

Request Instances Wizard [Cancel]

CHOOSE AN AMI | INSTANCE DETAILS | **CREATE KEY PAIR** | CONFIGURE FIREWALL | REVIEW

Public/private key pairs allow you to securely connect to your instance after it launches. To create a key pair, enter a name and click **Create & Download your Key Pair**. You will then be prompted to save the private key to your computer. Note, you only need to generate a key pair once - not each time you want to deploy an Amazon EC2 instance.

☒ Choose from your existing Key Pairs

Your existing Key Pairs*: IM2011

☐ Create a new Key Pair

☐ Proceed without a Key Pair

< Back | Continue >

AWS provisions an instance and returns user credentials

AWS

Elastic Beanstalk

Amazon

S3

Amazon

EC2

Amazon

VPC

Amazon

CloudWatch

Amazon

Elastic MapReduce

Amazon

CloudFront

AWS

CloudFormation

Amazon

RDS

Amazon

SNS

AWS

IAM

Navigation

Region:
US East (Virginia)

EC2 Dashboard

INSTANCES

Instances

Spot Requests

Reserved Instances

IMAGES

AMIs

Bundle Tasks

ELASTIC BLOCK STORE

Volumes

Snapshots

NETWORKING & SECURITY

Security Groups

Elastic IPs

Placement Groups

Load Balancers

Key Pairs

My Instances

Launch Instance

Instance Actions

Show/Hide

Refresh

Help

Viewing: All Instances

All Instance Types

1 to 3 of 3 Instances

	Name	Instance	AMI ID	Root Device	Type	Status	Security Groups	Key Pair Name	Monitoring	Virtualization	Placement
<input type="checkbox"/>	empty	i-a1c318cf	ami-e4a3578d	ebs	t1.micro	running	IM2001	IM2011	basic	paravirtual	
<input type="checkbox"/>	MyFirstInstance	i-3b7aa155	ami-76f0061f	ebs	m1.small	terminated	default		basic	paravirtual	
<input checked="" type="checkbox"/>		i-6176ad0f	ami-e4a3578d	ebs	t1.micro	running	IM2001	IM2011	detailed	paravirtual	

1 EC2 Instance selected

EC2 Instance: i-6176ad0f

Description

Monitoring

Tags

AMI: sles-11-sp1-v1.00.x86_64 (ami-e4a3578d)

Security Groups: IM2001

Status: running

VPC ID: -

Source/Dest. Check:

Placement Group:

RAM Disk ID: -

Key Pair Name: IM2011

Monitoring: detailed

Elastic IP: -

Root Device Type: ebs

Lifecycle: normal

Block Devices: sda1

Public DNS: ec2-50-16-69-93.compute-1.amazonaws.com

Private DNS: ip-10-196-229-93.ec2.internal

Private IP Address: 10.196.229.93

Launch Time: 2011-05-26 10:45 EDT

State Transition Reason:

Zone: us-east-1c

Type: t1.micro

Owner: 026317314573

Subnet ID: -

Virtualization: paravirtual

Reservation: r-ddfe6db1

Platform: -

Kernel ID: aki-427d952b

AMI Launch Index: 0

Root Device: sda1

Tenancy: default

AWS Services Catalog

Compute

EC2
Lightsail
Lambda
Batch
Elastic Beanstalk
Serverless Application Repository
AWS Outposts
EC2 Image Builder
AWS App Runner

Containers

Elastic Container Registry
Elastic Container Service
Elastic Kubernetes Service
Red Hat OpenShift Service on AWS

Storage

S3
EFS
FSx
S3 Glacier
Storage Gateway
AWS Backup
AWS Elastic Disaster Recovery

Database

RDS
ElastiCache
Neptune
Amazon QLDB
Amazon DocumentDB
Amazon Keyspaces
Amazon Timestream
DynamoDB
Amazon MemoryDB for Redis

Migration & Transfer

AWS Migration Hub
AWS Application Migration Service
Application Discovery Service
Database Migration Service
AWS Transfer Family
AWS Snow Family
DataSync
AWS Mainframe Modernization

Networking & Content Delivery

VPC
CloudFront
Route 53
API Gateway
Direct Connect
AWS App Mesh
AWS Cloud Map
Global Accelerator
Amazon VPC IP Address Manager
AWS Private 5G

Developer Tools

CodeStar
CodeCommit
CodeArtifact
CodeBuild
CodeDeploy
CodePipeline
Cloud9
CloudShell
X-Ray
AWS FIS

Customer Enablement

AWS IQ
Managed Services
Activate for Startups
Support

Robotics

AWS RoboMaker

Blockchain

Amazon Managed Blockchain

Satellite

Ground Station

Quantum Technologies

Amazon Braket

Management & Governance

AWS Organizations
CloudWatch
AWS Auto Scaling
CloudFormation
Config
OpsWorks
Service Catalog
Systems Manager
AWS AppConfig
Trusted Advisor
Control Tower
AWS License Manager
AWS Well-Architected Tool
AWS Health Dashboard
AWS Chatbot
Launch Wizard
AWS Compute Optimizer
Resource Groups & Tag Editor
Amazon Grafana
Amazon Prometheus
AWS Proton
AWS Resilience Hub
Incident Manager
CloudTrail

Media Services

Kinesis Video Streams
MediaConnect
MediaConvert

Machine Learning

Amazon SageMaker
Amazon Augmented AI
Amazon CodeGuru
Amazon DevOps Guru
Amazon Comprehend
Amazon Forecast
Amazon Fraud Detector
Amazon Kendra
Amazon Personalize
Amazon Polly
Amazon Rekognition
Amazon Textract
Amazon Transcribe
Amazon Translate
AWS DeepComposer
AWS DeepLens
AWS DeepRacer
AWS Panorama
Amazon Monitron
Amazon HealthLake
Amazon Lookout for Vision
Amazon Lookout for Equipment
Amazon Lookout for Metrics
Amazon Comprehend Medical
Amazon Lex

Analytics

Athena
Amazon Redshift
EMR
CloudSearch
Amazon OpenSearch Service
Kinesis
QuickSight
Data Pipeline
AWS Data Exchange
AWS Glue
AWS Lake Formation
MSK
AWS Glue DataBrew
Amazon FinSpace

Security, Identity, & Compliance

IAM
Resource Access Manager
Cognito
Secrets Manager
GuardDuty
Inspector
Amazon Macie
IAM Identity Center (successor to AWS Single Sign-On)
Certificate Manager
Key Management Service
CloudHSM
Directory Service
WAF & Shield
AWS Firewall Manager
Artifact
Security Hub
Detective

AWS Cost Management

AWS Cost Explorer
AWS Budgets
AWS Marketplace Subscriptions
AWS Application Cost Profiler
AWS Billing Conductor

Front-end Web & Mobile

AWS Amplify
AWS AppSync
Device Farm
Amazon Location Service

AR & VR

Amazon Sumerian

Application Integration

Step Functions
Amazon AppFlow
Amazon EventBridge
Amazon MQ
Simple Notification Service
Simple Queue Service
SWF
Managed Apache Airflow

Business Applications

Amazon Connect
Amazon Pinpoint
Amazon Honeycode
Amazon Chime
Amazon Simple Email Service
Amazon WorkDocs
Amazon WorkMail
Alexa for Business

End User Computing

WorkSpaces
AppStream 2.0
WorkSpaces Web

Internet of Things

IoT Core
FreeRTOS
IoT 1-Click
IoT Analytics
IoT Device Defender
IoT Device Management
IoT Events
IoT Greengrass
IoT SiteWise
IoT RoboRunner
IoT TwinMaker
AWS IoT FleetWise

Game Development

Amazon GameLift
Amazon GameSparks

How to build the backend using resources from cloud

Let's see how to leverage on-demand and elastic resources from cloud to build an extreme scale backend platform for a given application

- Let's first build a webapp
- Next let's create that webapp using resources from cloud
- Next we will progress towards building an extremely large scale application backend
-

Demos/Videos/Links

- Static website using AWS S3 service:
<https://docs.aws.amazon.com/AmazonS3/latest/userguide/HostingWebsiteOnS3Setup.html>
-

Next Week

- Reading List

- GFS: The Google File System

- <http://static.googleusercontent.com/media/research.google.com/en/us/archive/gfs-sosp2003.pdf>

- Mini Homework

- Sign up for AWS account. Sign up for AWS EC2 and S3 services.
 - Create a micro instance with Amazon Linux stack with appropriate keys and access control.
 - SSH into the instance you created. Take a screenshot and submit it. You submit in the courseworks under miniHW1 link in Assignments.

Lecture 2: IaaS Cloud and Cloud Programming

- API and CLI based access
 - <http://docs.aws.amazon.com/cli/latest/userguide/cli-chap-welcome.html>
- AWS access using Java SDK and CLI
- Learn how to use EC2 and S3 as example services
- Breaking down the steps - how AWS provided on-demand resource
- Create a web server and deploy your web application using AWS
- How to use on-demand infrastructure for regular applications?

Amazon EC2 Programming

- Amazon EC2 SDK for java on Eclipse
 - <http://aws.amazon.com/eclipse/>
 - <https://docs.aws.amazon.com/toolkit-for-eclipse/v1/user-guide/setup-install.html>
- AWS SDK with Python: <https://aws.amazon.com/sdk-for-python/>
- RESTful APIs for invoking EC2 APIs from Java

Deconstructing Amazon EC2 request machine API

- User goes to Amazon EC2 portal and specifies desired parameters for a machine
 - Resource: CPU, mem, disk
 - Stack: OS and possibly with additional software
- Amazon AWS Cloud manager (resource pool manager) provisions the user request
 - Finds appropriate physical resource
 - Dispatches the request to virtualization manager on the identified resource
 - Cloud Manager invokes EC2 API to provisions the request
- Virtualization manager on physical server
 - Copies the pre-built software stack (virtual appliance)
 - Provisions a guest VM and configures parameters (IP address, access rules,...) at run/boot time
- Cloud manager returns login credentials to user

1. User requests a machine with a desired Software stack, access rules

2. Cloud manager processes request

3. Identifies physical server where to instantiate

4. Virtualization mgr on the server launches a VM, copies virtual appliance and boots the VM with appropriate run-time configuration

5. Login credentials for user

6. User is provided instance details

