



AWS Service → Usage Hint Cheat Sheet

Context: Dining Concierge / Real-time Restaurant Recommendation System
(Assignment 1 → Q1 → Q2 → Q8 progression)

Service	Used For / Typical Question Context	When to Use It (Hint)	Example Phrase Trigger in Question
Amazon Lex	Chatbot frontend → understands user text (e.g., <i>"Find Italian food near me"</i>)	Use whenever question mentions chat interface, NLP, conversation, or intent extraction .	"User speaks to bot / chat interface" → Lex for intent + slots
AWS Lambda	Core logic processing unit for everything (API handlers, aggregators, background jobs)	Always use for stateless compute – POST handlers, rating aggregators, trending fetchers, scheduled jobs.	"Function to process data / trigger by event / schedule" → Lambda
API Gateway	Unified entry point for REST + WebSocket APIs	Always front end your Lambdas with API Gateway for auth, throttling, and CORS.	"Expose API / GET endpoint / handle HTTP requests" → API Gateway
Amazon SQS	Decouples producers and consumers; buffers requests for async processing	Use when you see asynchronous, event-driven, avoid	"Prevent bottlenecks / decouple flows" → SQS

overload, batch processing, queue.

Amazon SNS	Sends notifications / emails / mobile alerts to users	Mention when question adds “notify users / real-time alerts / publish-subscribe”.	“Send update to user / publish message” → SNS (or SES for email)
Amazon SES	Email delivery of recommendations or reports	Use for sending recommendation emails or system alerts .	“Send email / confirmation / daily recommendations” → SES
Amazon Cognito	Handles authentication and authorization with JWT tokens	Include whenever question mentions login, user auth, profiles, secure endpoints .	“User sign-in / token / auth required” → Cognito
Amazon DynamoDB	Primary data store for users, restaurants, likes, ratings, trending	Always when you need fast, scalable NoSQL reads/writes < 10 ms.	“Store user / rating / trending data / aggregate” → DynamoDB
Amazon OpenSearch (ElasticSearch)	Full-text and geo-spatial search engine for queries	Use when question mentions search, filter, geo-based, ranking by distance or rating .	“Find nearby restaurants / search by keyword” → OpenSearch
Amazon S3 + CloudFront	Frontend hosting, asset delivery, CDN caching for static site	Always present at top of architecture. Mention for frontend UI / CDN / low-latency delivery .	“Static website / UI / frontend content” → S3 + CloudFront

ElastiCache (Redis)	In-memory cache for hot data (trending, ratings, nearby restaurants)	Use for low-latency (<1 ms) re-reads of popular data / avoid DB hot keys.	“Cache / hot data / reduce latency / store trending” → Redis
Amazon CloudWatch	Monitoring + Metrics + Triggers (Lambda logs, schedules)	Include for metrics, alarms, scheduled tasks, monitoring .	“Every 5 min / daily job / monitor errors / logs” → CloudWatch
Amazon EventBridge	Scheduler & event bus for periodic Lambda jobs	Mention when you see ‘run every 15 min / 2 AM daily’ / trigger Lambda on event .	“Scheduled job / event-based workflow” → EventBridge
AWS WAF + Shield + KMS	Security layer (attack protection + encryption)	Add whenever you discuss security / data encryption / DDoS protection .	“Secure data / protect API / encrypt PII” → WAF + KMS
AWS X-Ray	Distributed tracing and latency debugging	Optional for questions on latency bottlenecks / monitoring .	“Trace requests / identify slow Lambda” → X-Ray

How to Think During the Exam

If the question asks for... → instantly map to these:

**Question Theme /
Keyword**

Use These Services

Why

Chatbot / NLP	Lex → Lambda → SQS	Conversational intent handling pipeline
User Feedback / Ratings / Likes	API Gateway → Lambda → DynamoDB → SQS → Aggregator Lambda → DynamoDB update	Async writes + aggregation without blocking
Trending / Popularity / Realtime Data	External API → Lambda → DynamoDB TrendingCache → Redis → OpenSearch	Cache hot data, reduce latency
Search / Nearby Results	OpenSearch + DynamoDB + Redis	Combine geo search with cached aggregates
Wait-Time or Offers (API Integration)	External API → Lambda → DynamoDB Cache + TTL + EventBridge schedule	Cached third-party data to save API costs
Email / Alerts	SES / SNS + SQS → Lambda	Asynchronous notifications to users
Auth / Profiles / Users	Cognito + DynamoDB	Manage users securely with JWT
Scalability / Burst Traffic	SQS + Lambda auto-scale + Redis cache + DynamoDB on-demand	Decouple + cache + auto-scale
Monitoring / Automation	CloudWatch + EventBridge + X-Ray	Scheduled tasks + performance visibility



Quick Example Mapping

Question Type	Core Services Involved
Q1 – Personalized Recommendations	Lex, API Gateway, Lambda, SQS, DynamoDB (UserProfiles, UserLikes, TrendingCache), OpenSearch, SES
Q2 – Wait-Times & Offers	Q1 + Lambda Aggregator + EventBridge + Extra DynamoDB Caches + External APIs
Q8 – Ratings & Trending	Q1 + Q2 + ElastiCache (Redis) + DynamoDB (Ratings, TrendingCache) + External Trending API + Multi-Region Scaling
Assignment 1 (Core)	Lex → Lambda → SQS → DynamoDB → SES (+ API Gateway + Cognito)



Exam Hints (Memory Triggers)

If Question Mentions...	Immediately Think of...
“Chatbot”	Lex + Lambda
“High Traffic / Async / Queue”	SQS
“Schedule / every 5 minutes”	EventBridge or CloudWatch Event

"Trending / Cache / Hot Data"	ElastiCache (Redis)
"Personalized / User Based"	Cognito + UserProfiles + OpenSearch
"Search Nearby / Geo Queries"	OpenSearch (geo_point)
"Rating / Feedback / Write Ops"	Lambda + DynamoDB Streams + Aggregator
"Real-Time Notification / Alert"	SNS / WebSocket / SES
"Authentication / Token"	Cognito
"Monitoring / Logs / Metrics"	CloudWatch + X-Ray
"Secure Endpoints"	WAF + KMS
