

# System design from sample paper

## Enhanced Personalized Restaurant Recommendation Engine

### 1. Data Stores

To build an enhanced personalized recommendation engine for the Dining Concierge chatbot, several data stores are integrated to handle various aspects of data management, personalization, and real-time adaptability. Below is a detailed list and description of each data store, including their schemas and indexing mechanisms.

#### a. User Profile Store (Amazon DynamoDB)

- **Purpose:** Stores user-specific data, including preferences, past interactions, and feedback.
- **Data Stored:**
  - **UserID** (Primary Key): Unique identifier for each user.
  - **Preferences:** JSON object containing preferred cuisines, locations, dining times, etc.
  - **Past Searches:** Array of search queries with timestamps.
  - **Feedback:** Array of liked/disliked restaurants with timestamps.
- **Schema Example:**

json

Copy code

```
{ "UserID": "user123", "Preferences": { "cuisines": ["Japanese", "Italian"], "locations": ["Manhattan", "Brooklyn"], "diningTime": "7 pm" }, "PastSearches": [ { "query": "Japanese in Manhattan", "timestamp": "2024-10-25T19:00:00Z" } ], "Feedback": [ { "restaurantID": "rest001", "liked": true, "timestamp": "2024-10-25T20:00:00Z" } ] }
```

- **Indexing Mechanisms:**
  - **Primary Key:** UserID ensures unique and fast retrieval of user profiles.
  - **Global Secondary Index (GSI):** Index on Preferences.cuisines to facilitate quick querying based on cuisine preferences.

#### b. Restaurant Data Store (Amazon DynamoDB & Amazon Elasticsearch Service)

- **Purpose:** Stores comprehensive restaurant information and facilitates efficient searching and filtering.

##### i. DynamoDB Table: yelp-restaurants

- **Data Stored:**
  - **BusinessID** (Primary Key): Unique identifier for each restaurant.
  - **Name:** Name of the restaurant.

- **Address:** Physical address.
- **Coordinates:** Geolocation data (latitude and longitude).
- **NumberOfReviews:** Total number of reviews.
- **Rating:** Average rating.
- **ZipCode:** Postal code.
- **InsertedAtTimestamp:** Timestamp of data insertion.

- **Schema Example:**

json

Copy code

```
{ "BusinessID": "rest001", "Name": "Sushi Nakazawa", "Address": "23 Commerce St", "Coordinates": {"lat": 40.71427, "lon": -74.00597}, "NumberOfReviews": 1500, "Rating": 4.8, "ZipCode": "10014", "InsertedAtTimestamp": "2024-04-01T12:00:00Z" }
```

- **Indexing Mechanisms:**
  - **Primary Key:** BusinessID for unique identification.
  - **GSI on Cuisine:** To enable quick retrieval based on cuisine types.

## ii. Elasticsearch Index: restaurants

- **Data Stored:**
  - **RestaurantID:** Same as BusinessID in DynamoDB.
  - **Cuisine:** Type of cuisine offered.
- **Schema Example:**

json

Copy code

```
{ "RestaurantID": "rest001", "Cuisine": "Japanese" }
```

- **Indexing Mechanisms:**
  - **Inverted Index** on Cuisine to facilitate full-text search and quick filtering based on cuisine types.
  - **Geo Indexing** on Coordinates for efficient geographical proximity searches.

## c. Recommendation Queue (Amazon SQS)

- **Purpose:** Acts as a buffer for recommendation requests to ensure asynchronous processing.
- **Data Stored:**
  - Messages containing user requests with relevant parameters (e.g., UserID, Cuisine, Location).

- **Schema Example:**

json

Copy code

```
{ "UserID": "user123", "Cuisine": "Japanese", "Location": "Manhattan", "NumberOfPeople": 2, "DiningTime": "7 pm", "Email": "user@example.com" }
```

- **Indexing Mechanism:**
  - SQS does not require traditional indexing; it handles message ordering and visibility internally.

#### d. Trending Data Store (Amazon DynamoDB)

- **Purpose:** Stores trending restaurants data based on user interactions.
- **Data Stored:**
  - **RestaurantID:** Unique identifier.
  - **TrendScore:** Calculated based on the number of likes.
  - **LastUpdated:** Timestamp of the last update.
- **Schema Example:**

json

Copy code

```
{ "RestaurantID": "rest001", "TrendScore": 150, "LastUpdated": "2024-10-31T10:00:00Z" }
```

- **Indexing Mechanisms:**
  - **Primary Key:** RestaurantID.
  - **GSI on TrendScore:** To enable quick retrieval of top trending restaurants.

#### e. State Management Store (Amazon DynamoDB)

- **Purpose:** Maintains the state of user interactions for personalized experiences.
- **Data Stored:**
  - **UserID (Primary Key):** Unique identifier.
  - **LastSearch:** Details of the last search performed.
- **Schema Example:**

json

Copy code

```
{ "UserID": "user123", "LastSearch": { "Cuisine": "Japanese", "Location": "Manhattan" } }
```

- **Indexing Mechanisms:**
  - **Primary Key:** UserID for quick state retrieval.

## 2. APIs

To support the enhanced functionalities, several new APIs are integrated into the system. These APIs facilitate personalized recommendations, trending data retrieval, and state management.

### a. New APIs

#### 1. GetPersonalizedRecommendations API

- **Endpoint:** /recommendations/personalized
- **Method:** GET
- **Description:** Fetches personalized restaurant recommendations based on user preferences and past interactions.
- **Parameters:** UserID, optional filters (e.g., location, cuisine).

#### 2. GetTrendingRecommendations API

- **Endpoint:** /recommendations/trending
- **Method:** GET
- **Description:** Retrieves trending restaurant recommendations based on collective user interactions.
- **Parameters:** Location, optional filters.

#### 3. SubmitFeedback API

- **Endpoint:** /feedback/submit
- **Method:** POST
- **Description:** Allows users to submit feedback (like/dislike) on recommended restaurants.
- **Payload:** UserID, RestaurantID, FeedbackType.

#### 4. GetLastSearch API

- **Endpoint:** /user/last-search
- **Method:** GET
- **Description:** Retrieves the last search parameters of a user for automatic recommendations.
- **Parameters:** UserID.

### b. Low-Level Backend Design

To handle high traffic efficiently, the backend is designed using AWS services that support scalability, event-driven architecture, and asynchronous processing.

#### 1. API Gateway

- **Role:** Acts as the entry point for all API requests.

- **Features:**
  - **Throttling and Rate Limiting:** Ensures API stability under high traffic.
  - **Authentication:** Secures APIs using AWS Cognito or API keys.

## 2. AWS Lambda Functions

- **Functionality:**
  - **PersonalizedRecommendationsHandler:** Processes requests for personalized recommendations.
  - **TrendingRecommendationsHandler:** Processes requests for trending recommendations.
  - **FeedbackHandler:** Handles user feedback submissions.
  - **LastSearchHandler:** Retrieves the last search state for a user.
- **Scalability:** Automatically scales based on incoming request volume.
- **Statelessness:** Ensures that each invocation is independent, enhancing reliability.

## 3. Amazon DynamoDB

- **Role:** Serves as the primary data store for user profiles, state management, and trending data.
- **Features:**
  - **DAX (DynamoDB Accelerator):** Optional in-memory caching for read-heavy operations.
  - **Auto Scaling:** Automatically adjusts read/write capacity based on traffic.

## 4. Amazon Elasticsearch Service

- **Role:** Facilitates efficient searching and filtering based on restaurant attributes.
- **Features:**
  - **Shard and Replica Configuration:** Ensures high availability and performance.
  - **Kibana Integration:** For monitoring and visualizing search queries.

## 5. Amazon SQS

- **Role:** Manages asynchronous processing of recommendation requests.
- **Features:**
  - **FIFO Queues:** Ensures ordered processing if required.
  - **Dead-Letter Queues:** Handles failed message processing.

## 6. Amazon SNS (Simple Notification Service)

- **Role:** Notifies other services or triggers workflows based on specific events (e.g., new feedback submission).

## 7. Amazon SES (Simple Email Service)

- **Role:** Sends personalized recommendation emails to users.

## 8. Amazon EventBridge (formerly CloudWatch Events)

- **Role:** Manages scheduled tasks and event-driven triggers.

### c. Scalability and Resilience Features

- **Auto Scaling:** Both API Gateway and Lambda automatically scale to handle varying loads.
- **Load Balancing:** API Gateway efficiently distributes incoming requests.
- **Retry Mechanisms:** Implemented in SQS and Lambda to handle transient failures.
- **Monitoring and Logging:** Amazon CloudWatch provides metrics and logs for all services, enabling proactive scaling and issue resolution.

### 3. System Design Architecture (High-Level Backend Design)

Below is a high-level architecture diagram illustrating the integration of the personalized recommendation engine and dynamic data system with the existing chatbot.

*Note: As this is a text-based response, please visualize the architecture as described below.*

#### Architecture Components and Flow

##### 1. User Interaction:

- Users interact with the frontend hosted on **Amazon S3**, which communicates with the backend via **API Gateway**.

##### 2. API Gateway:

- Routes requests to appropriate **AWS Lambda** functions:
  - **PersonalizedRecommendationsHandler**
  - **TrendingRecommendationsHandler**
  - **FeedbackHandler**
  - **LastSearchHandler**

##### 3. Lambda Functions:

- **PersonalizedRecommendationsHandler:**
  - Retrieves user preferences and past interactions from **DynamoDB**.
  - Queries **Amazon Elasticsearch Service** for matching restaurants.
  - Fetches detailed information from **DynamoDB**.
  - Compiles and sends recommendations via **Amazon SES**.
- **TrendingRecommendationsHandler:**
  - Queries **Trending Data Store** in **DynamoDB**.

- Retrieves trending restaurants from **Elasticsearch** and **DynamoDB**.
- Sends recommendations via **Amazon SES**.
- **FeedbackHandler**:
  - Processes user feedback and updates **Trending Data Store**.
- **LastSearchHandler**:
  - Retrieves the last search parameters for a user to offer automatic recommendations.

#### 4. Data Stores:

- **DynamoDB**:
  - **User Profile Store**: Stores user-specific data.
  - **Restaurant Data Store**: Comprehensive restaurant information.
  - **Trending Data Store**: Tracks trending restaurants.
  - **State Management Store**: Maintains user interaction states.
- **Amazon Elasticsearch Service**:
  - Facilitates efficient search and filtering of restaurants based on cuisine and geographic proximity.

#### 5. Asynchronous Processing:

- **Amazon SQS**:
  - Manages recommendation request queues.
  - **Lambda Workers** (e.g., **LF2**) process queue messages to fetch and send recommendations.

#### 6. Real-Time Data Integration:

- **External Data Sources** (e.g., Yelp API) feed updates into **Elasticsearch** and **DynamoDB** via scheduled **Lambda** functions or **EventBridge** triggers.

#### 7. Notification Service:

- **Amazon SES** sends personalized recommendation emails to users based on processed data.

### 4. Feedback Loop, Real-Time Processing, and Adaptability

#### a. Feedback Loop

The system incorporates a robust feedback loop to continuously refine and personalize recommendations based on user interactions and preferences.

##### 1. User Feedback Collection:

- Users interact with the frontend and can "like" recommended restaurants.
- These interactions are sent to the **SubmitFeedback API**, which invokes the **FeedbackHandler Lambda**.

## 2. Feedback Processing:

- **FeedbackHandler Lambda** updates the **User Profile Store** in **DynamoDB** with the feedback.
- It also updates the **Trending Data Store** by incrementing the TrendScore for liked restaurants.

## 3. Recommendation Refinement:

- Future personalized recommendations consider updated user preferences and trending data.
- This ensures that recommendations evolve with user behavior and broader trends.

## b. Real-Time Data Handling and Adaptability

The system is designed to adapt to real-time data changes, such as restaurant availability or new openings, ensuring that recommendations remain relevant and up-to-date.

### 1. External Data Integration:

- **Data Source Assumption:** Assume integration with the **Yelp API** for real-time restaurant data updates.
- **Lambda Function:** Scheduled **Lambda** functions fetch updates from the Yelp API at regular intervals (e.g., every hour).

### 2. Data Update Pipeline:

- **Lambda** fetches new or updated restaurant data.
- Updates are pushed to both **DynamoDB** and **Elasticsearch**:
  - **DynamoDB:** Ensures comprehensive and up-to-date restaurant details.
  - **Elasticsearch:** Facilitates efficient searching and filtering based on the latest data.

### 3. Handling Restaurant Availability:

- **Lambda Workers** periodically check restaurant statuses (e.g., open/closed) via the Yelp API.
- Updates are reflected in the data stores to exclude unavailable restaurants from recommendations.

### 4. Dynamic Recommendations:

- **Lambda Handlers** dynamically query the latest data from **Elasticsearch** and **DynamoDB** to provide current recommendations.



- **Trending Recommendations** are recalculated based on the latest user interactions and feedback.

### c. AWS Components and Services for Real-Time Management

- **Amazon Kinesis (Optional):**
  - For handling high-throughput real-time data streams from external APIs.
- **AWS Lambda:**
  - Processes real-time data updates and user feedback.
- **Amazon EventBridge:**
  - Orchestrates event-driven workflows, triggering Lambda functions based on specific events or schedules.
- **Amazon DynamoDB Streams:**
  - Captures changes in DynamoDB tables for real-time processing and integration with other services.

## 5. Data Pipeline / Event Flow

The data pipeline ensures seamless flow from user interaction to the delivery of personalized recommendations, incorporating feedback and real-time data updates.

### a. Data Pipeline Steps

1. **User Interaction:**
  - User engages with the chatbot via the frontend hosted on **S3**.
  - Requests for recommendations are sent through **API Gateway** to the respective Lambda handlers.
2. **Recommendation Request Processing:**
  - **PersonalizedRecommendationsHandler** Lambda retrieves user data from **DynamoDB**.
  - Queries **Elasticsearch** for matching restaurants based on preferences and geographic proximity.
  - Fetches detailed restaurant information from **DynamoDB**.
  - Compiles a list of 5 personalized and 5 trending recommendations.
  - Sends the recommendations via **Amazon SES**.
3. **Feedback Submission:**
  - User "likes" a restaurant, triggering the **SubmitFeedback API**.
  - **FeedbackHandler Lambda** updates user profiles and trending data.
4. **Trending Data Update:**

- **FeedbackHandler** updates the **Trending Data Store** in **DynamoDB**.
  - The **TrendingRecommendationsHandler** Lambda uses this data to prioritize trending restaurants.
5. **Real-Time Data Updates:**
- **Scheduled Lambda Functions** fetch updates from the Yelp API.
  - Updates are pushed to **DynamoDB** and **Elasticsearch**.
  - Ensures that recommendations reflect the latest restaurant data.
6. **State Management:**
- **LastSearchHandler** Lambda retrieves the last search parameters.
  - Provides automatic recommendations based on past searches when the user returns.
7. **Asynchronous Processing:**
- **Amazon SQS** queues handle high-volume recommendation requests.
  - **Lambda Workers** process queue messages, ensuring scalability and reliability.

## **b. Event Flow Diagram**

1. **User Requests Recommendations:**
  - Frontend → API Gateway → PersonalizedRecommendationsHandler Lambda
2. **Lambda Processes Request:**
  - Fetch user data from DynamoDB
  - Query Elasticsearch for restaurants
  - Retrieve details from DynamoDB
  - Send email via SES
3. **User Provides Feedback:**
  - Frontend → SubmitFeedback API → FeedbackHandler Lambda
4. **FeedbackHandler Updates Data Stores:**
  - Update User Profile Store and Trending Data Store in DynamoDB
5. **Scheduled Data Updates:**
  - EventBridge triggers DataUpdate Lambda
  - Fetch and update data from Yelp API to DynamoDB and Elasticsearch
6. **Queue Processing:**
  - New recommendation requests are placed in SQS
  - Lambda Workers poll SQS and process messages

## Assumptions

- **External Data Source:** The Yelp API is used for fetching real-time restaurant data, including new openings and availability statuses.
- **User Authentication:** Users are authenticated via AWS Cognito to secure API endpoints and manage user identities.
- **Data Volume:** The system is designed to handle a large user base with high-frequency interactions and data updates.
- **Email Sending Limits:** Amazon SES is configured to handle the expected email volume, with appropriate sending limits and verified domains.

## Conclusion

The enhanced Personalized Restaurant Recommendation Engine leverages AWS services to deliver a dynamic, user-centric experience. By integrating personalized recommendations, trending data, and real-time adaptability, the system ensures that users receive relevant and up-to-date restaurant suggestions. The architecture emphasizes scalability, resilience, and efficient data management, making it robust enough to handle a growing user base and evolving data landscapes.