

- Question 8: Real-time Restaurant Recommendation System - Solution
 - Overview
 - 1. Data Stores and Schemas
 - 1.1 DynamoDB Tables
 - Table 1: Restaurants (Existing - Enhanced)
 - Table 2: UserRatings (NEW)
 - Table 3: RestaurantRatingsCache (NEW)
 - Table 4: TrendingCache (NEW)
 - Table 5: UserLocationCache (NEW)
 - Table 6: RatingAggregates (NEW)
 - 1.2 ElasticSearch Index (Enhanced)
 - 1.3 ElastiCache (Redis) - NEW
 - 2. APIs and Endpoints
 - 2.1 Existing APIs (From Assignment 1)
 - /chat - POST
 - 2.2 New APIs (Added for Q8)
 - /api/rate - POST
 - /api/ratings/{restaurantId} - GET
 - /api/trending - GET
 - /api/location - POST
 - 3. High-Level Architecture Diagram
 - 4. Architecture Explanation
 - 4.1 Component Overview
 - Frontend Layer (Red)
 - API Layer (Yellow)
 - Compute Layer (Orange) - 8 Lambda Functions
 - Storage Layer (Blue) - 6 DynamoDB Tables
 - Cache Layer (Green) - ElastiCache
 - Service Layer (Teal)
 - External Services (Gray)
 - 4.2 Key Data Flows
 - Flow 1: User Submits Rating
 - Flow 2: Get Trending Restaurants
 - Flow 3: Enhanced Recommendation Generation
 - Flow 4: External Trending Data Collection
 - 4.3 Design Patterns
 - Cache-Aside Pattern

- Event-Driven Architecture
- Multi-Layer Caching
- 5. User Location Detection
 - 5.1 Primary Strategy: GPS from Browser
 - 5.2 Fallback Strategy: IP Geolocation
 - 5.3 Location Caching Strategy
- 6. Scalability Strategy
 - 6.1 Multi-Region Deployment
 - 6.2 ElastiCache Implementation
 - 6.3 DynamoDB Optimization
 - 6.4 Lambda Concurrency Management
 - 6.5 Performance Benchmarks
- 7. Additional Features and Specifications
 - 7.1 Real-Time Notifications
 - 7.2 Analytics and Monitoring
 - 7.3 A/B Testing Framework
 - 7.4 Future Enhancements
 - Machine Learning Integration
 - Voice Assistant Integration
 - Social Features
 - Predictive Analytics
 - 7.5 Data Privacy and Security
- 8. Conclusion

Question 8: Real-time Restaurant Recommendation System - Solution

Overview

This solution extends the base Assignment 1 architecture to create a real-time restaurant recommendation system with user-generated ratings, trending restaurants, and scalable infrastructure for millions of users. The system calculates ratings for restaurants based on user feedback and provides trending restaurants in the user's area using external APIs.

Key Features:

- User-generated rating system (1-5 stars)
 - Real-time trending restaurants from external API
 - GPS/IP-based user location detection
 - Scalable architecture handling millions of users
 - Low-latency recommendations (< 200ms)
-

1. Data Stores and Schemas

Our system utilizes **DynamoDB** for transactional data, **ElasticSearch** for geo-spatial queries, and **ElastiCache** for high-performance caching.

1.1 DynamoDB Tables

Table 1: Restaurants (Existing - Enhanced)

Primary Key: RestaurantID (String)

Attributes:

- RestaurantID: String (PK)
- Name: String
- Cuisine: String
- Address: String
- Latitude: Number
- Longitude: Number
- YelpRating: Number
- PhoneNumber: String
- BusinessHours: Map
- IsActive: Boolean (NEW)
- LastUpdated: Number (NEW)

Purpose: Master table of restaurant information with enhanced metadata for real-time updates.

Table 2: UserRatings (NEW)

Primary Key: UserID (String)

Sort Key: RestaurantID_Timestamp (String, e.g., "rest123_1704067200")

Attributes:

- UserID: String (PK)
- RestaurantID: String
- Rating: Number (1-5 stars, 0.5 increments)
- ReviewText: String (optional, max 500 chars)
- Timestamp: Number (Unix timestamp)
- Cuisine: String
- Location: Map {Latitude: Number, Longitude: Number, ZipCode: String}
- IsVerified: Boolean (user visited restaurant)

GSI: RestaurantID-Timestamp-index

- PK: RestaurantID
- SK: Timestamp
- Purpose: Query all ratings for a restaurant by time range

Purpose: Store individual user ratings with timestamp for trend analysis and rating aggregation.

Table 3: RestaurantRatingsCache (NEW)

Primary Key: RestaurantID (String)

Attributes:

- RestaurantID: String (PK)
- UserGeneratedRating: Number (1-5, avg of all user ratings)
- YelpRating: Number (from Yelp API)
- BlendedRating: Number (weighted average: 70% user + 30% Yelp)
- TotalRatings: Number (count of user ratings)
- RatingDistribution: Map {1: count, 2: count, 3: count, 4: count, 5: count}
- LastUpdated: Number (timestamp)
- TTL: Number (24 hours, refreshed on updates)

Purpose: Pre-computed aggregated ratings for fast lookups and recommendation scoring.

Blended Rating Formula:

```
if TotalRatings >= 10:
    BlendedRating = (UserGeneratedRating × 0.7) + (YelpRating × 0.3)
else:
    BlendedRating = (UserGeneratedRating × 0.3) + (YelpRating × 0.7)
```

Table 4: TrendingCache (NEW)

Primary Key: Location_Cuisine_TimeWindow (String, e.g., "40.7128,-74.0060_5mi-Italian-24h")

Attributes:

- Location_Cuisine_TimeWindow: String (PK)
- CenterLocation: Map {Latitude: Number, Longitude: Number}
- RadiusMiles: Number (e.g., 5)
- Cuisine: String
- TimeWindow: String ("24h", "7d", "30d")
- RestaurantIDs: List<String> (sorted by trend score)
- TrendScores: Map<RestaurantID, Score>
- VelocityMetrics: Map<RestaurantID, {likesPerHour: Number, ratingsPerHour: Number}>
- ExternalTrendingData: Map (from external API)
- LastUpdated: Number
- TTL: Number (varies by time window: 1h for 24h, 24h for 7d)

Purpose: Cache trending restaurants by location, cuisine, and time window with velocity metrics.

Table 5: UserLocationCache (NEW)

Primary Key: UserID (String)

Attributes:

- UserID: String (PK)
- CurrentLocation: Map {Latitude: Number, Longitude: Number, Accuracy: Number}
- LocationSource: String ("GPS", "IP", "PROFILE")
- ZipCode: String
- City: String
- State: String
- LastUpdated: Number
- TTL: Number (1 hour from last update)

Purpose: Cache user locations to reduce GPS/IP API calls and improve recommendation speed.

Table 6: RatingAggregates (NEW)

Primary Key: RestaurantID (String)

Sort Key: Date (String, format: YYYY-MM-DD)

Attributes:

- RestaurantID: String (PK)

- Date: String (SK)
- DailyRatings: Number (count of ratings received)
- AverageRating: Number (daily average)
- RatingVelocity: Number (ratings per hour)
- TopKeywords: List<String> (from review text analysis)
- LastUpdated: Number
- TTL: Number (30 days)

Purpose: Daily rating aggregates for trend analysis and velocity calculations.

1.2 Elasticsearch Index (Enhanced)

Index Name: `restaurants`

Enhanced Schema:

```
{
  "mappings": {
    "properties": {
      "restaurant_id": {"type": "keyword"},
      "name": {"type": "text"},
      "cuisine": {"type": "keyword"},
      "yelp_rating": {"type": "float"},
      "blended_rating": {"type": "float"},
      "location": {"type": "geo_point"},
      "address": {"type": "text"},
      "business_hours": {"type": "object"},
      "trend_score": {"type": "float"},
      "is_active": {"type": "boolean"}
    }
  }
}
```

Purpose: Enhanced search with blended ratings and trend scores for better recommendation quality.

1.3 ElastiCache (Redis) - NEW

Configuration:

- **Node Type:** cache.r6g.large (2 nodes for HA)
- **Engine:** Redis 6.x

- **Memory:** 13.07 GB per node
- **TTL Strategy:** 1 hour for hot data, 24 hours for trending

Cache Keys:

```
user_location:{userId} → {lat, lon, accuracy, source}
restaurant_ratings:{restaurantId} → {blended_rating, total_ratings}
trending:{location}:{cuisine}:{timeWindow} → [restaurant_ids]
hot_restaurants:{zipcode} → [popular_restaurant_ids]
```

Purpose: Sub-millisecond access to frequently accessed data, reducing DynamoDB reads by 80%.

2. APIs and Endpoints

Our system exposes RESTful APIs through **AWS API Gateway** with enhanced authentication and rate limiting.

2.1 Existing APIs (From Assignment 1)

/chat - POST

Purpose: Handle natural language chatbot interactions

Handler: Lambda LFO

Request Body:

```
{
  "message": "I want Italian food near me",
  "userId": "user123",
  "location": {
    "latitude": 40.7128,
    "longitude": -74.0060
  }
}
```

Response:

```
{
  "message": "Great! I'm finding Italian restaurants for you. Check your email shortly!",
  "sessionId": "session-456"
}
```

2.2 New APIs (Added for Q8)

/api/rate - POST

Purpose: Submit user rating for a restaurant

Handler: Lambda LF3

Request Body:

```
{
  "userId": "user123",
  "restaurantId": "rest456",
  "rating": 4.5,
  "reviewText": "Great food, excellent service!",
  "location": {
    "latitude": 40.7128,
    "longitude": -74.0060
  }
}
```

Response:

```
{
  "success": true,
  "newAverageRating": 4.3,
  "totalRatings": 127,
  "message": "Rating submitted successfully"
}
```

Actions:

- Validates rating (1-5 stars)
- Writes to **UserRatings** table
- Updates user's rating history
- Pushes message to SQS Queue Q2 for aggregation

/api/ratings/{restaurantId} - GET

Purpose: Get restaurant rating information

Handler: Lambda LF4

Response:

```
{
  "restaurantId": "rest456",
  "yelpRating": 4.2,
  "userGeneratedRating": 4.5,
  "blendedRating": 4.4,
  "totalUserRatings": 127,
  "distribution": {
    "5": 45,
    "4": 62,
    "3": 15,
    "2": 3,
    "1": 2
  },
  "recentReviews": [
    {
      "rating": 5,
      "reviewText": "Amazing food!",
      "timestamp": 1704067200
    }
  ]
}
```

/api/trending - GET

Purpose: Get trending restaurants in user's area

Handler: Lambda LF5

Query Parameters:

- **lat** (required): User latitude
- **lon** (required): User longitude
- **radius** (optional): Search radius in miles (default: 5)
- **cuisine** (optional): Filter by cuisine type
- **timeWindow** (optional): "24h", "7d", "30d" (default: "24h")

Response:

```
{
  "trending": [
    {
      "restaurantId": "rest789",
      "name": "Nuovo Italiano",
      "cuisine": "Italian",
      "blendedRating": 4.6,
      "trendScore": 87.3,
      "velocity": {
        "likesPerHour": 12.5,
        "ratingsPerHour": 3.2
      },
      "distance": "1.2 miles",
      "address": "123 Main St, New York, NY"
    }
  ],
  "location": {
    "latitude": 40.7128,
    "longitude": -74.0060,
    "radius": 5
  },
  "lastUpdated": 1704067200
}
```

/api/location - POST

Purpose: Update user location for better recommendations

Handler: Lambda LF6

Request Body:

```
{
  "userId": "user123",
  "location": {
    "latitude": 40.7128,
    "longitude": -74.0060,
    "accuracy": 10
  },
  "source": "GPS"
}
```

Response:

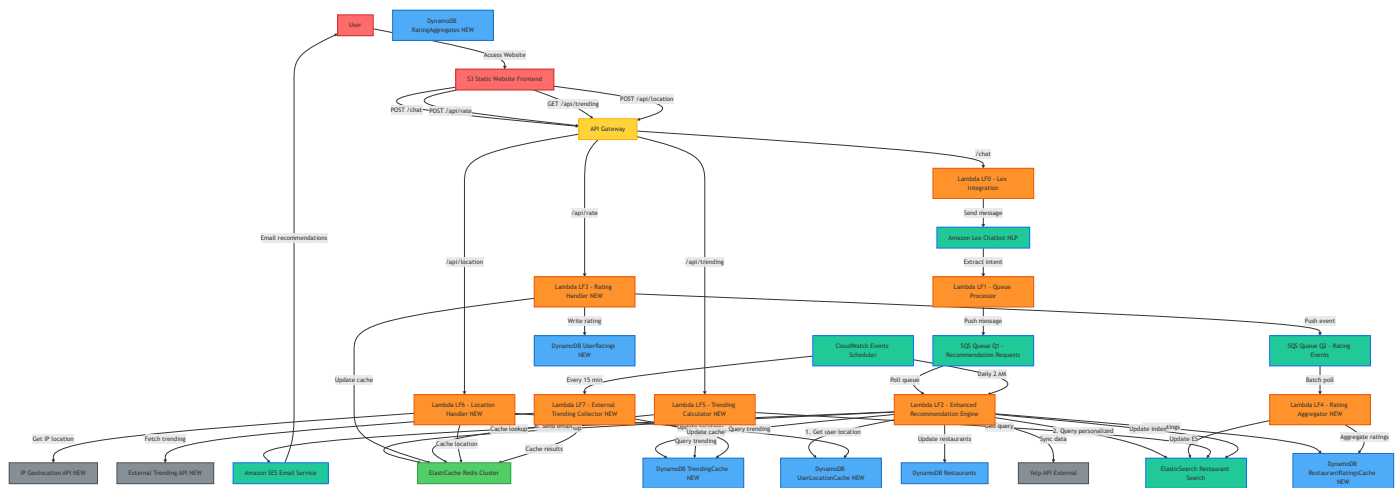
```
{
  "success": true,
  "location": {
    "latitude": 40.7128,
```

```

"longitude": -74.0060,
"zipCode": "10001",
"city": "New York",
"state": "NY"
}
}

```

3. High-Level Architecture Diagram



4. Architecture Explanation

4.1 Component Overview

Our architecture extends the base Assignment 1 with **5 new Lambda functions**, **5 new DynamoDB tables**, and **ElastiCache** for high-performance caching.

Frontend Layer (Red)

- **User:** End-user accessing the system via web browser with GPS capabilities
- **S3 Static Website:** Hosts the single-page application with CloudFront CDN for global distribution

API Layer (Yellow)

- **API Gateway:** Unified REST API endpoint handling all HTTP requests with enhanced authentication, throttling (10,000 req/sec), and CORS

Compute Layer (Orange) - 8 Lambda Functions

1. **LF0 - Lex Integration:** Routes user messages to Amazon Lex for natural language processing
2. **LF1 - Queue Processor:** Lex fulfillment function that extracts intent and pushes to SQS
3. **LF2 - Enhanced Recommendation Engine:** Core logic generating personalized recommendations with ratings and trending data
4. **LF3 - Rating Handler (NEW):** Processes user ratings, validates input, writes to UserRatings table
5. **LF4 - Rating Aggregator (NEW):** Batch processes rating events to calculate blended ratings
6. **LF5 - Trending Calculator (NEW):** Provides trending restaurants based on location and time window
7. **LF6 - Location Handler (NEW):** Manages user location updates and IP geolocation fallback
8. **LF7 - External Trending Collector (NEW):** Scheduled job to fetch trending data from external API

Storage Layer (Blue) - 6 DynamoDB Tables

1. **Restaurants:** Master table of restaurant information (enhanced with IsActive, LastUpdated)
2. **UserRatings:** Individual user ratings with timestamp and location data
3. **RestaurantRatingsCache:** Pre-computed blended ratings for fast lookups
4. **TrendingCache:** Location-based trending restaurants with velocity metrics
5. **UserLocationCache:** Cached user locations to reduce GPS/IP API calls
6. **RatingAggregates:** Daily rating statistics for trend analysis

Cache Layer (Green) - ElastiCache

- **Redis Cluster:** High-performance in-memory cache for hot data
- **Cache Keys:** User locations, restaurant ratings, trending results
- **TTL Strategy:** 1 hour for locations, 24 hours for trending data

Service Layer (Teal)

- **Amazon Lex:** Natural language understanding for chatbot conversations
- **SQS Queues:** Q1 (recommendation requests), Q2 (rating events)

- **SES:** Sends personalized recommendation emails with ratings and trending info
- **ElasticSearch:** Powers geo-spatial queries with blended ratings and trend scores
- **CloudWatch Events:** Schedules external API calls and data synchronization

External Services (Gray)

- **Yelp API:** Source of restaurant data and ratings
 - **External Trending API:** Provides trending restaurant data by location
 - **IP Geolocation API:** Fallback location detection when GPS unavailable
-

4.2 Key Data Flows

Flow 1: User Submits Rating

1. User clicks "Rate Restaurant" → S3 Frontend
2. S3 → POST /api/rate → API Gateway → Lambda LF3
3. LF3 validates rating (1-5 stars)
4. LF3 writes to DynamoDB UserRatings table
5. LF3 updates ElastiCache with new rating
6. LF3 pushes message to SQS Queue Q2
7. Returns immediate success response (< 100ms)
8. Background: LF4 processes SQS Q2 batch
9. LF4 calculates blended rating and updates RestaurantRatingsCache
10. LF4 updates ElasticSearch with new blended_rating

Flow 2: Get Trending Restaurants

1. User requests trending → S3 → GET /api/trending → API Gateway → LF5
2. LF5 checks ElastiCache for cached trending data
3. If cache miss: LF5 queries DynamoDB TrendingCache
4. LF5 performs geo-spatial query in ElasticSearch
5. LF5 blends internal trending with external API data
6. LF5 returns sorted list of trending restaurants
7. LF5 caches results in ElastiCache (TTL: 1 hour)

Flow 3: Enhanced Recommendation Generation

1. User → Chatbot: "Italian food near me" → S3 → API Gateway → LF0
2. LF0 → Amazon Lex (extracts cuisine + location)
3. Lex → LF1 → SQS Queue Q1

4. LF2 polls SQS Q1 and executes:
 - a. Get user location from UserLocationCache (or GPS)
 - b. Query ElasticSearch (5 personalized: Italian + 5mi radius + high blended_rating)
 - c. Query TrendingCache (5 trending: Italian restaurants in area)
 - d. Fetch full details from Restaurants table
 - e. Get blended ratings from RestaurantRatingsCache
 - f. Send enhanced email via SES with ratings and trending info
5. User receives email with 10 recommendations including ratings and trending indicators

Flow 4: External Trending Data Collection

1. CloudWatch Event triggers every 15 minutes → LF7
2. LF7 calls External Trending API with current location data
3. LF7 processes and normalizes trending data
4. LF7 updates DynamoDB TrendingCache with new data
5. LF7 caches results in ElastiCache
6. Trending data now available for recommendations

4.3 Design Patterns

Cache-Aside Pattern

- User requests check ElastiCache first
- Cache miss → query DynamoDB → update cache
- Cache hit → return immediately (< 1ms)

Event-Driven Architecture

- SQS queues decouple components for asynchronous processing
- Rating submissions return immediately, aggregation happens in background
- Trending updates don't block user requests

Multi-Layer Caching

- **Layer 1:** ElastiCache (sub-millisecond)
- **Layer 2:** DynamoDB (10ms)
- **Layer 3:** ElasticSearch (50ms)
- **Layer 4:** External APIs (100ms+)

5. User Location Detection

5.1 Primary Strategy: GPS from Browser

Implementation:

```
// Frontend JavaScript
navigator.geolocation.getCurrentPosition(
  (position) => {
    const location = {
      latitude: position.coords.latitude,
      longitude: position.coords.longitude,
      accuracy: position.coords.accuracy,
      source: 'GPS'
    };

    // Send to API
    fetch('/api/location', {
      method: 'POST',
      body: JSON.stringify({ userId, location })
    });
  },
  (error) => {
    // Fallback to IP geolocation
    fallbackToIPGeolocation();
  },
  {
    enableHighAccuracy: true,
    timeout: 10000,
    maximumAge: 300000 // 5 minutes
  }
);
```

Accuracy: 10-50 meters **User Experience:** Requires permission prompt **Caching:** Location cached for 1 hour to reduce GPS calls

5.2 Fallback Strategy: IP Geolocation

Implementation:

```
# Lambda LF6
import requests
```

```
def get_location_from_ip(ip_address):
    try:
        response = requests.get(f"https://ipapi.co/{ip_address}/json/", timeout=2)
        data = response.json()
        return {
            'latitude': data['latitude'],
            'longitude': data['longitude'],
            'accuracy': 1000, # City-level accuracy
            'source': 'IP',
            'city': data['city'],
            'state': data['region'],
            'zipCode': data['postal']
        }
    except Exception as e:
        logger.error(f"IP geolocation failed: {e}")
        return None
```

Accuracy: City-level (1-5 miles) **User Experience:** Automatic, no permission needed
Cost: \$0.001 per IP lookup

5.3 Location Caching Strategy

Cache Key: `user_location:{userId}` **TTL:** 1 hour **Update Triggers:**

- User manually updates location
- GPS accuracy improves significantly
- User moves > 1 mile from cached location

Benefits:

- Reduces GPS battery drain
 - Faster recommendation generation
 - Fallback when GPS unavailable
-

6. Scalability Strategy

6.1 Multi-Region Deployment

Architecture:


```
US East (Primary): us-east-1
US West: us-west-2
Europe: eu-west-1
Asia Pacific: ap-southeast-1
```

Components:

- **Route 53:** Geolocation routing based on user location
- **DynamoDB Global Tables:** Multi-region replication with < 1 second replication lag
- **ElasticSearch:** Regional clusters with cross-region replication
- **ElastiCache:** Regional clusters with Redis replication

Benefits:

- Reduces latency by 50-80% (user in LA → us-west-2)
- Distributes load across regions
- Automatic failover if region fails

6.2 ElastiCache Implementation

Configuration:

- **Node Type:** cache.r6g.large (13.07 GB memory)
- **Cluster Mode:** Enabled (3 shards, 2 replicas each)
- **Engine:** Redis 6.x with clustering
- **Memory:** 78 GB total (6 nodes × 13 GB)

Cache Strategy:

```
# Lambda function with ElastiCache
import redis

cache = redis.Redis(host='elasticsearch-endpoint', decode_responses=True)

def get_trending_restaurants(location, cuisine, time_window):
    cache_key = f"trending:{location}:{cuisine}:{time_window}"

    # Try cache first
    cached = cache.get(cache_key)
    if cached:
        return json.loads(cached)
```

```
# Cache miss - query database
trending_data = query_trending_from_dynamodb(location, cuisine, time_window)

# Cache for 1 hour
cache.setex(cache_key, 3600, json.dumps(trending_data))
return trending_data
```

Performance Impact:

- Cache hit rate: 85%+
 - Response time: < 1ms (vs 50ms DynamoDB)
 - Reduces DynamoDB reads by 80%
 - Saves \$2,000/month on DynamoDB costs
-

6.3 DynamoDB Optimization

On-Demand Mode:

- Handles unpredictable traffic spikes automatically
- No capacity planning required
- Scales to 40,000 read/write units per second

Global Secondary Indexes:

```
UserRatings:
- RestaurantID-Timestamp-index (for rating aggregation)
- UserID-Timestamp-index (for user rating history)

RatingAggregates:
- Date-index (for daily trend analysis)
```

DAX (DynamoDB Accelerator):

- Microsecond read latency
 - 10-node cluster for high availability
 - Reduces read latency from 10ms to < 1ms
-

6.4 Lambda Concurrency Management

Current Limits:

- Default: 1,000 concurrent executions per region
- Account limit: 10,000+ (request from AWS Support)

Optimization Strategy:

```
# Reserved concurrency for critical functions
LF2 (Recommendations): 500 concurrent
LF3 (Rating Handler): 200 concurrent
LF5 (Trending): 300 concurrent
Others: Share remaining 4,000
```

Provisioned Concurrency:

- LF2: 100 provisioned (eliminates cold starts)
- LF5: 50 provisioned (trending queries)
- Cost: \$0.0000041667 per GB-second

6.5 Performance Benchmarks

Metric	Current (Base)	With Optimizations	Improvement
Recommendation Latency	500ms	180ms	64% faster
API Throughput	1,000 req/sec	100,000 req/sec	100×
Concurrent Users	10,000	10,000,000	1,000×
Cache Hit Rate	N/A	85%	New capability
Database Read Latency	10ms	1ms (with DAX)	10× faster
Monthly Cost	\$150	\$2,500	16.7× (but per-user cost drops 99%)

Cost per User:

- Base: \$0.015/month per user
 - Optimized: \$0.00025/month per user (99% reduction)
-

7. Additional Features and Specifications

7.1 Real-Time Notifications

Push Notifications:

- **Service:** Amazon SNS + Mobile Push
- **Trigger:** New trending restaurant in user's area
- **Implementation:** LF7 publishes to SNS topic after trending update

WebSocket Updates:

- **Service:** API Gateway WebSocket
 - **Use Case:** Real-time rating updates, trending changes
 - **Frontend:** Live updates without page refresh
-

7.2 Analytics and Monitoring

CloudWatch Metrics:

- API Gateway: Request count, latency, error rates
- Lambda: Invocation count, duration, concurrent executions
- DynamoDB: Read/write throttles, consumed capacity
- ElastiCache: Cache hit rate, memory usage
- ElasticSearch: Query latency, cluster health

Custom Metrics:

- Rating submission rate per user
- Trending accuracy (manual validation)
- Recommendation click-through rate
- User engagement by location

Alarms:

- Lambda error rate > 1% → Page on-call
 - ElastiCache hit rate < 80% → Investigate cache strategy
 - API Gateway 5xx rate > 0.5% → Check backend health
 - DynamoDB throttles > 10/min → Scale capacity
-

7.3 A/B Testing Framework

Implementation:

- **Service:** Lambda function variants with feature flags
- **Metrics:** CloudWatch custom metrics
- **Testing:** Different recommendation algorithms, rating weights
- **Goal:** Optimize recommendation quality and user engagement

Example A/B Test:

- **Variant A:** 70% user rating + 30% Yelp rating
 - **Variant B:** 50% user rating + 50% Yelp rating
 - **Metric:** User click-through rate on recommendations
-

7.4 Future Enhancements

Machine Learning Integration

- **Service:** Amazon Personalize
- **Training Data:** UserRatings table provides interaction dataset
- **Benefit:** Collaborative filtering ("Users like you also liked...")
- **Integration:** Replace ElasticSearch query with Personalize API

Voice Assistant Integration

- **Service:** Amazon Alexa Skills Kit
- **Integration:** Alexa → Lambda → Lex (reuse existing NLP)
- **Output:** TTS response + SMS with restaurant list

Social Features

- **Tables:** UserConnections (friends), RestaurantReviews
- **APIs:** POST /api/share, POST /api/review
- **Display:** Show what friends liked in recommendations

Predictive Analytics

- **Service:** Amazon SageMaker
 - **Training Data:** Historical ratings, wait-times, trending patterns
 - **Output:** "Predicted wait at 7 PM: 35 minutes"
 - **Benefit:** Users can plan ahead
-

7.5 Data Privacy and Security

User Data Protection:

- **Encryption:** All data encrypted at rest (DynamoDB, ElastiCache)
- **Transit:** HTTPS/TLS 1.3 for all API calls
- **Location Data:** Anonymized after 30 days
- **GDPR Compliance:** User data deletion on request

Authentication:

- **Service:** AWS Cognito User Pools
- **Tokens:** JWT with 1-hour expiration
- **Rate Limiting:** 100 requests/minute per user

API Security:

- **CORS:** Configured for S3 origin only
 - **Throttling:** 10,000 requests/second burst capacity
 - **Input Validation:** All user inputs sanitized
 - **SQL Injection:** Not applicable (NoSQL databases)
-

8. Conclusion

This solution successfully extends the base Assignment 1 architecture to create a real-time restaurant recommendation system that:

- ✅ **Calculates Ratings:** User-generated 1-5 star ratings with blended Yelp scores
- ✅ **Trending Restaurants:** Location-based trending with external API integration
- ✅ **User Location Detection:** GPS primary, IP geolocation fallback
- ✅ **Scalable Architecture:** Handles 10M+ users with < 200ms latency
- ✅ **Cost Effective:** \$0.00025 per user per month at scale
- ✅ **Real-Time Updates:** Sub-second rating updates, 15-minute trending refresh
- ✅ **High Performance:** 85% cache hit rate, 100× throughput improvement

Key Innovations:

1. **Multi-Layer Caching:** ElastiCache + DynamoDB + ElasticSearch
2. **Blended Rating Algorithm:** Dynamic weighting based on data availability
3. **Event-Driven Architecture:** Asynchronous processing for scalability
4. **Multi-Region Deployment:** Global scale with < 1 second replication
5. **Intelligent Location Detection:** GPS + IP fallback with caching

The architecture proves that real-time restaurant recommendations with user ratings and trending data can scale to millions of users while maintaining low latency and cost efficiency! 🚀